

240-229: SDA (Operating Systems session)

Lecture 2: Process and Thread Concepts

Associate Professor Dr. Sangsuree Vasupongayya

sangsuree.v@psu.ac.th

Department of computer Engineering

Prince of Songkla University

Outline

□ Lecture:

- Process Concept
- Operations on Processes
- Thread Concept

□ Labs:


- Process
- Thread

Thinking

What is a process?

Process – a program in execution; process execution must progress in sequential fashion



What a process need during execution?



What does OS need to know about a process?

Thinking

Are all processes in the system having a different state during their lifetime?



How many state?

How do I know what state a process is at?

Process Concept

- **Process** – a program in execution; process execution must progress in sequential fashion
- A process includes:
 - program counter
 - stack
 - data section

heap: is a memory that is dynamically allocated during process runtime






Diagram of Process State



Process Control Block (PCB)


Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information




Process Scheduling Queues

- Job queue – set of all processes in the system
- Ready queue – set of all processes residing in main memory, ready and waiting to execute
- Device queues – set of processes waiting for an I/O device
- Processes migrate among the various queues



Thinking

What happen when you have more than one process on the system at the same time?




CPU: only a limited # of core

I/O: only one monitor


Memory: only limited space

Context Switch

- ☐ When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- ☐ Context-switch time is overhead; the system does no useful work while switching
- ☐ Time dependent on hardware support



CPU Switch From Process to Process




Process Creation

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program



Process Creation (fork)

- Fork creates a new process by copying the calling process
- The new process has its own
 - Memory address space (copied from parent)
 - Instructions (same program as parent!)
 - Data
 - Stack
 - Register set (copied from parent)
 - Process table entry in the OS

Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**)
 - Output data from child to parent (via **wait**)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - Some operating system do not allow child to continue if its parent terminates
 - All children terminated - *cascading termination*

Lab: Fork

```
int main(int argc, char *argv[])
{
    int temp = 10;
    pid_t pid;
    /* fork a child process */
    pid = fork(); // return 0 to the child process
                  // whereas the nonzero process id of
                  // the child is returned to the parent

    if(pid < 0) { /* error occurred */
        fprintf(stderr,"Fork Failed");
        exit(-1);
    }
    else if(pid == 0) { /* child process */
        // sleep(5);
        printf("This is a child process. pid = %d, parent pid = %d\n",getpid(),getppid());
        sleep(2);
        printf("This is a child process. I woke up! \n");
        execvp("/bin/ls","ls",NULL);
        //while(1);
        temp++;
        printf("CHILD: temp = %d\n",temp);
    }
    else { /* parent process*/
        //sleep(5);
        printf("This is a parent process. pid = %d, child pid = %d\n",getpid(),pid);
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("PARENT: Child Complete.\n");
        printf("PARENT: temp = %d\n",temp);
        exit(0);
    }
}
```

Process have the following components

- An address space
- A collection of operating system state
- A CPU context .. Or thread of control


What is a Thread?

- A thread executes a stream of instructions
- A thread is an abstraction for control-flow


Thread is a processor context and stack

- Allocated a CPU by a scheduler
- Executes in a memory address space

Single and Multithreaded Processes





single-threaded process




multithreaded process

3 processes vs 3 threads




Environment (resource)

Example: Word processor



- One thread can wait while the other can still process

3 threads



Why use threads?


- Responsiveness
 - Handle asynchronous events
- Economic
 - Blocking due to I/O
 - Overlap computation
- Resource sharing
 - Low cost communication via shared memory
- Utilization of Multiprocessor Architectures

Sequential vs. concurrent threads


- Sequential threads




- Concurrent threads running simultaneously on two processors



- Concurrent threads (with gaps in their executions) interleaved on one processor




A Web Server Example




A Web Server Example

Why is this not a good web server design?




A Web Server Example

Fork more process




A Web Server Example

Create a thread



A Web Server Example

Have a pool of threads



Common Thread Strategies

□ Manager/worker


- Manager thread handles I/O
- Manager assigns work to worker threads
- Worker threads created dynamically
- ... or allocated from a *thread-pool*

□ Pipeline

- Each thread handles a different stage of an assembly line
- Threads hand work off to each other in a *producer-consumer* relationship


Kernel-Level Threads

Thread-switching code is in the kernel



User-Level Threads Package

The thread-switching code is in user space



Thread Libraries


A thread library provides the programmer with an API for creating and managing threads

- Provide a library entirely in user space with no kernel support
- Implement a kernel-level library supported directly by the operating system
 - Code and data structures for the library exist in kernel space
 - Invoking a function in API for the library results in a system call to the kernel

pthread

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)
- First thread in main creates the others
 - `pthread_create()`
 - `pthread_join()`
 - `pthread_exit()`

Using Create, Join and Exit



Java Threads

- Java threads are managed by the JVM
- Java threads may be created by:
 - Implementing the Runnable interface

```
public interface Runnable  
{  
    public abstract void run();  
}
```

