

240-229: SDA (Operating Systems session)

Lecture 3: Concurrency

Associate Professor Dr.Sangsuree Vasupongayya
sangsuree.v@psu.ac.th
Department of Computer Engineering
Prince of Songkla University

Outline

□ Lecture:

- Interprocess communication
- Concurrency concept
- The Critical-Section Problem
- Solutions to Critical-section
- Deadlock
- Classic Problems of Synchronization

□ Lab:


- Pipe
- Shared memory
- Mutex lock

Interprocess Communication (IPC)


- Process can be independent or cooperating
 - Independent – is not affecting or is not affected by the other processes executing in the system
 - Cooperating – can affect or be affected by the other processes executing in the system
- Why IPC?
 - Cooperating process use IPC to exchange data and information
- 2 Types
 - Message passing (e.g., pipe)
 - Shared memory

Interprocess Communication


Message Passing



Shared Memory




IPC supported by Windows



- ☐ Clipboard: very loosely coupled exchange medium, on agreed data format


gRPC (RPC = Remote Procedure Call)



- <http://www.grpc.io/>
- a high performance, open-source universal RPC framework
- Stubby used by Google as a single general-purpose RPC infrastructure to connect the large number of microservices running within and across multiple data centers.
- based around the idea of defining a service, specifying the methods that can be called remotely with their parameters and return types
- Supports various popular programming languages such as C
- uses protocol buffers as the Interface Definition Language (IDL)

Concurrency

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers
 - having an integer **count** that keeps track of the number of full buffers



consumer-producer problem

- One of the solutions to a consumer-producer problem that fills **all** the buffers
 - having an integer **count** that keeps track of the number of full buffers
 - Initially, count is set to 0
 - count is incremented by the producer after it produces a new buffer
 - Count is decremented by the consumer after it consumes a buffer

Producer

- Produce an item
- Add the item to the buffer
- If the buffer is full, the producer does not add any item

```
while (count == BUFFER_SIZE)
    ; // do nothing

    // add an item to the buffer
    ++count;
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
```

Consumer

- Remove an item from the buffer
- If the buffer is empty, the consumer does not remove any item

```
while (count == 0)
    ; // do nothing

// remove an item from the buffer
--count;
item = buffer[out];
out = (out + 1) % BUFFER_SIZE;
```

Producer-Consumer

```
while (count == BUFFER_SIZE)
    ; // do nothing

// add an item to the buffer
++count;
buffer[in] = item;
in = (in + 1) % BUFFER_SIZE;
```

Producer

Consumer

```
while (count == 0)
    ; // do nothing

// remove an item from the buffer
--count;
item = buffer[out];
out = (out + 1) % BUFFER_SIZE;
```

**Let count = 5
And BUFFER_SIZE = 5**

Producer-Consumer

```
while (count == BUFFER_SIZE)
    ; // do nothing

// add an item to the buffer
++count;
buffer[in] = item;
in = (in + 1) % BUFFER_SIZE;

Producer

Consumer

while (count == 0)
    ; // do nothing

// remove an item from the buffer
--count;
item = buffer[out];
out = (out + 1) % BUFFER_SIZE;
```

Let count = 5
And BUFFER_SIZE = 5

Case 1:
producer
 loop do nothing
consumer
 remove 1 item
producer
 add 1 item

Producer-Consumer

```
while (count == BUFFER_SIZE)
    ; // do nothing

// add an item to the buffer
++count;
buffer[in] = item;
in = (in + 1) % BUFFER_SIZE;

Producer

Consumer

while (count == 0)
    ; // do nothing

// remove an item from the buffer
--count;
item = buffer[out];
out = (out + 1) % BUFFER_SIZE;
```

**Let count = 5
And BUFFER_SIZE = 5**

Case 1:
producer
loop do nothing
consumer
remove 1 item
producer
add 1 item

Case 2:
consumer
remove 1 item
producer
add 1 item

Both cases, count = 5

Producer-Consumer

```
while (count == BUFFER_SIZE)
    ; // do nothing

// add an item to the buffer
++count;
buffer[in] = item;
in = (in + 1) % BUFFER_SIZE;

Producer

Consumer

while (count == 0)
    ; // do nothing

// remove an item from the buffer
--count;
item = buffer[out];
out = (out + 1) % BUFFER_SIZE;
```

**Let count = 5
And BUFFER_SIZE = 5**

Case 1:
producer
loop do nothing
consumer
remove 1 item
producer
add 1 item

Case 2:
consumer
remove 1 item
producer
add 1 item

Both cases, count = 5

Race Condition

- ❑ `count++` could be implemented as

```
register1 = count  
register1 = register1 + 1  
count = register1
```

- ❑ `count--` could be implemented as

```
register2 = count  
register2 = register2 - 1  
count = register2
```

- ❑ Consider this execution interleaving with “`count = 5`” initially:

S0: producer executes `register1 = count`

S1: producer executes `register1 = register1 + 1`

S2: consumer executes `register2 = count`

S3: consumer executes `register2 = register2 - 1`

S4: producer executes `count = register1`

S5: consumer executes `count = register2`

Race Condition (cont.)

Consider this execution interleaving with “count = 5” initially:

- S0: producer executes `register1 = count` {reg1 = 5}
- S1: producer executes `register1 = register1 + 1` {reg1 = 6}
- S2: consumer executes `register2 = count` {reg2 = 5}
- S3: consumer executes `register2 = register2 - 1` {reg2 = 4}
- S4: producer executes `count = register1` {count = 6 }
- S5: consumer executes `count = register2` {count = 4}

Critical Section

Section of code where shared data is accessed which is a segment of code that the process may be changing common variables, updating a table, writing a file, and so on

- Race Condition - When there is concurrent access to shared data and the final outcome depends upon order of execution
- Solution:
 - Only one process at a time can execute in its critical section

Solution to a Critical-Section Problem

```
while (true) {  
    entry section  
    critical section  
    exit section  
    remainder section  
}
```

- Entry Section - Code that requests permission to enter its critical section.
- Critical Section
- Exit Section - Code that is run after exiting the critical section
- Remainder Section

Solution to a Critical-Section Problem

Mutual Exclusion - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections

Progress - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely

Bounded Waiting - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted

- Assume that each process executes at a nonzero speed
- No assumption concerning relative speed of the N processes

Critical-Section using Locks

```
while (true) {
```

acquire lock

critical section

release lock

remainder section

```
}
```

Semaphore

- a synchronization tool
- Semaphore S – integer variable
- Two standard operations modify
 - S: acquire()
 - R: release()
- Less complicated
- Can only be accessed via two indivisible (atomic) operations

```
acquire() {  
    while value <= 0  
        ; // no-op  
    value--;  
}  
  
release() {  
    value++;  
}
```

Semaphore as General Synchronization Tool

- Counting semaphore – integer value can range over an unrestricted domain
- Binary semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - Also known as mutex locks

```
Semaphore S = new Semaphore();  
  
S.acquire();  
  
    // critical section  
  
S.release();  
  
    // remainder section
```

Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue.
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

Implementation with no Busy waiting (Cont.)

- Implementation of **acquire()**:

```
acquire(){
    value--;
    if (value < 0) {
        add this process to list
        block;
    }
}
```

- Implementation of **release()**:

```
release(){
    value++;
    if (value <= 0) {
        remove a process P from list
        wakeup(P);
    }
}
```

Problems with Semaphores

- Incorrectly use can cause timing errors
 - release() acquire()
 - release() ... release()
 - acquire() ... acquire()
 - acquire() or release() alone
- Errors are difficult to detect

Error only happens if some particular execution sequences take place and these sequences do not always occur
- Correct use of semaphore operations:
 - acquire() release()


Pthreads Synchronization

- Pthreads API is OS-independent
- It provides:
 - mutex locks
 - condition variables
- Non-portable extensions include:
 - read-write locks
 - spin locks

Atomic Transactions

- Assures that operations happen as a single logical unit of work, in its entirety, or not at all
- Related to field of database systems
- Challenge is assuring atomicity despite computer system failures
- Transaction - collection of instructions or operations that performs single logical function
 - Transaction is series of **read** and **write** operations
 - Terminated by **commit** (transaction successful) or **abort** (transaction failed) operation
 - Aborted transaction must be **rolled back** to undo any changes it performed


Deadlock and Starvation



Deadlock – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes




Starvation – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.



The Deadlock Problem

A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.



- Traffic only in one direction
- Each section of a bridge can be viewed as a resource
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
- Several cars may have to be backed up if a deadlock occurs
- Starvation is possible

Deadlock Characterization


Deadlock can arise if four conditions hold simultaneously

- **Mutual exclusion:** only one process at a time can use a resource.
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain (e.g., $P_0 \rightarrow P_1 \rightarrow P_2 \rightarrow \dots \rightarrow P_n \rightarrow P_0$)


Resource-Allocation Graph

A set of vertices V and a set of edges E .


- Process




- Resource Type with 4 instances





- P_i requests instance of R_j




- P_i is holding an instance of R_j



Resource Allocation Graph




Resource Allocation Graph (cont.)



Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

Recovery from Deadlock: Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- In which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer to completion.
 - Resources the process has used.
 - Resources process needs to complete.
 - How many processes will need to be terminated.
 - Is process interactive or batch?

Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.

Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem

Bounded-Buffer Problem

- A.k.a consumer-producer problem
 - Two processes (the producer and the consumer) who share a common buffer.
 - The producer's job is to generate a piece of data, put it into the buffer and start again.
 - The consumer's job is to consume the data (i.e., remove it from the buffer) one piece at a time.
 - The problem
 - The producer won't try to add data into the buffer if it's full
 - The consumer won't try to remove from an empty buffer

Bounded-Buffer Solution

□ Solution

- N buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value N .

Bounded-Buffer Solution

□ Solution

- N buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value N .

```
/* producer */
while (true) {
    empty.acquire();
    mutex.acquire();

    // add an item

    mutex.release();
    full.release();
}

/* consumer */
while (true) {
    full.acquire();
    mutex.acquire();

    // remove an item

    mutex.release();
    empty.release();
}
```

Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do NOT perform any update
 - Writers – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.

Readers-Writers Solution

- Semaphore **mutex** initialized to 1.
- Semaphore **db** initialized to 1.
- Integer **readerCount** initialized to 0.

Readers-Writers Solution

- Semaphore **mutex** initialized to 1.
- Semaphore **db** initialized to 1.
- Integer **readerCount** initialized to 0.

```
/* reader */
while (true) {

    db.acquire();

    // read from the database

    db.release();

}
```

Readers-Writers Solution

- Semaphore **mutex** initialized to 1.
- Semaphore **db** initialized to 1.
- Integer **readerCount** initialized to 0.

```
/* reader */
while (true) {

    ++readerCount;
    db.acquire();

    // read from the database

    --readerCount;
    db.release();

}
```

Readers-Writers Solution

- Semaphore **mutex** initialized to 1.
- Semaphore **db** initialized to 1.
- Integer **readerCount** initialized to 0.

```
/* reader */
while (true) {
    mutex.acquire();
    ++readerCount;
    if(readerCount == 1) { db.acquire(); } // first reader
    mutex.release();

    // read from the database

    mutex.acquire();
    --readerCount;
    if(readerCount == 0) { db.release(); } // last reader
    mutex.release();
}
```

Readers-Writers Solution

```
/* writer */
while (true) {
    db.acquire();
    // update the database
    db.release();
}

/* reader */
while (true) {
    mutex.acquire();
    ++readerCount;
    if(readerCount == 1) { db.acquire(); } // first reader
    mutex.release();
    // read from the database
    mutex.acquire();
    --readerCount;
    if(readerCount == 0) { db.release(); } // last reader
    mutex.release();
}

acquire() {
    while value <= 0
        ; // no-op
    value--;
}

release() {
    value++;
}
```