

# 240-229: SDA (Operating Systems session)

---

## Lecture 5: Memory Management

Associate Professor Dr. Sangsuree Vasupongayya

[sangsuree.v@psu.ac.th](mailto:sangsuree.v@psu.ac.th)

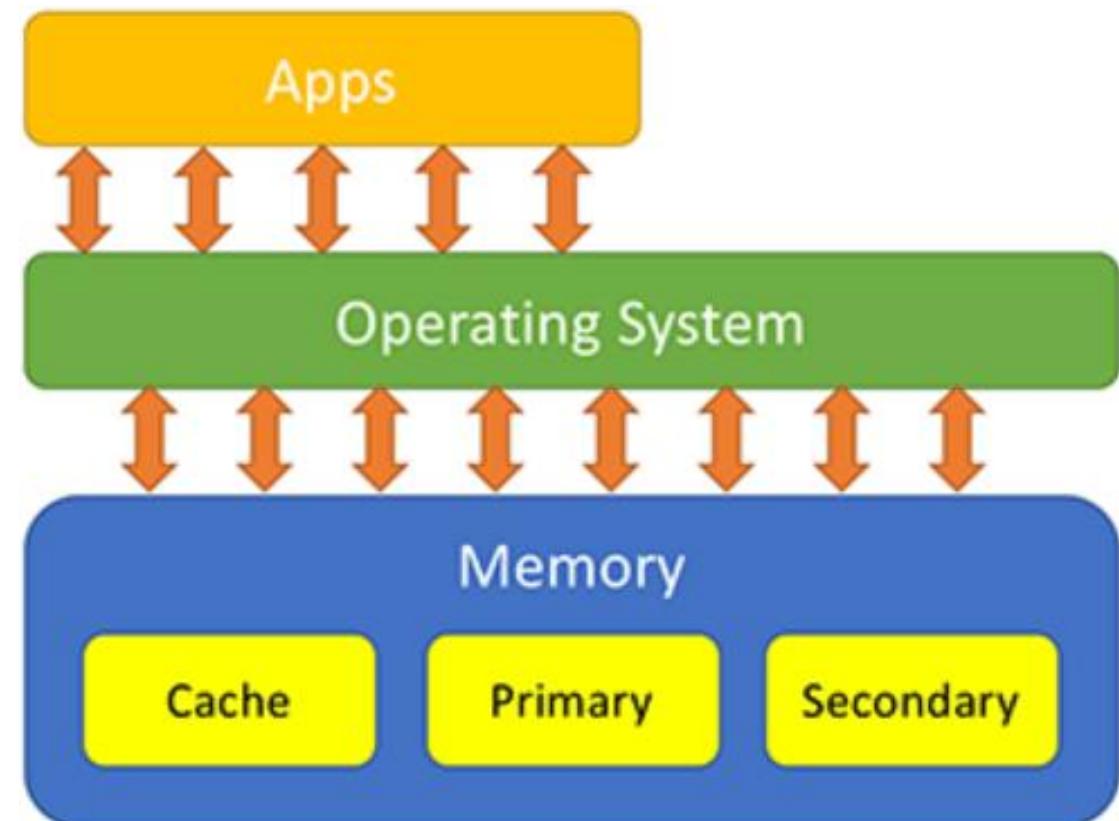
Department of Computer Engineering

Prince of Songkla University

# Outline

---

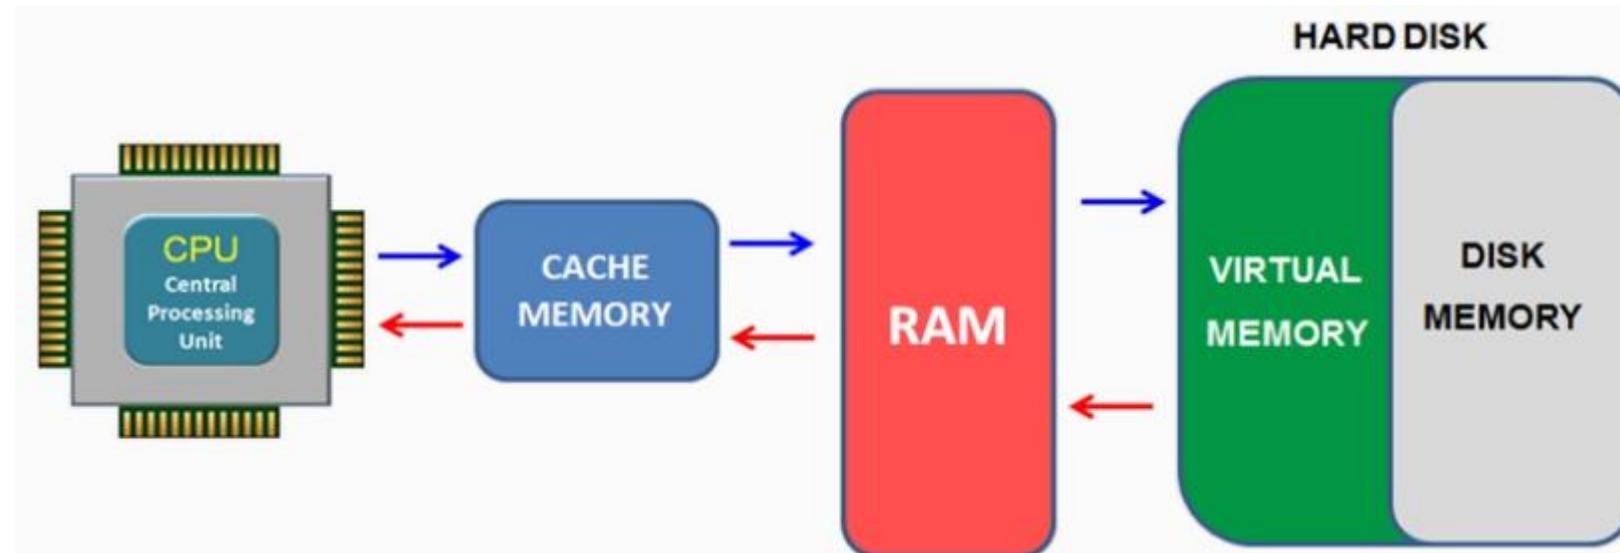
- Background
- Memory Management Requirements
- Paging
- Segmentation
- Virtual memory
- Replacement policy
- Kernel memory allocation



# Background

---

- Program must be brought (from disk) into memory for it to be executed
  - Main memory and registers are the only storage that the CPU can access directly
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure the correctness of operations



# Managing main memory

---

## Objective

- Allow as many processes to be executed

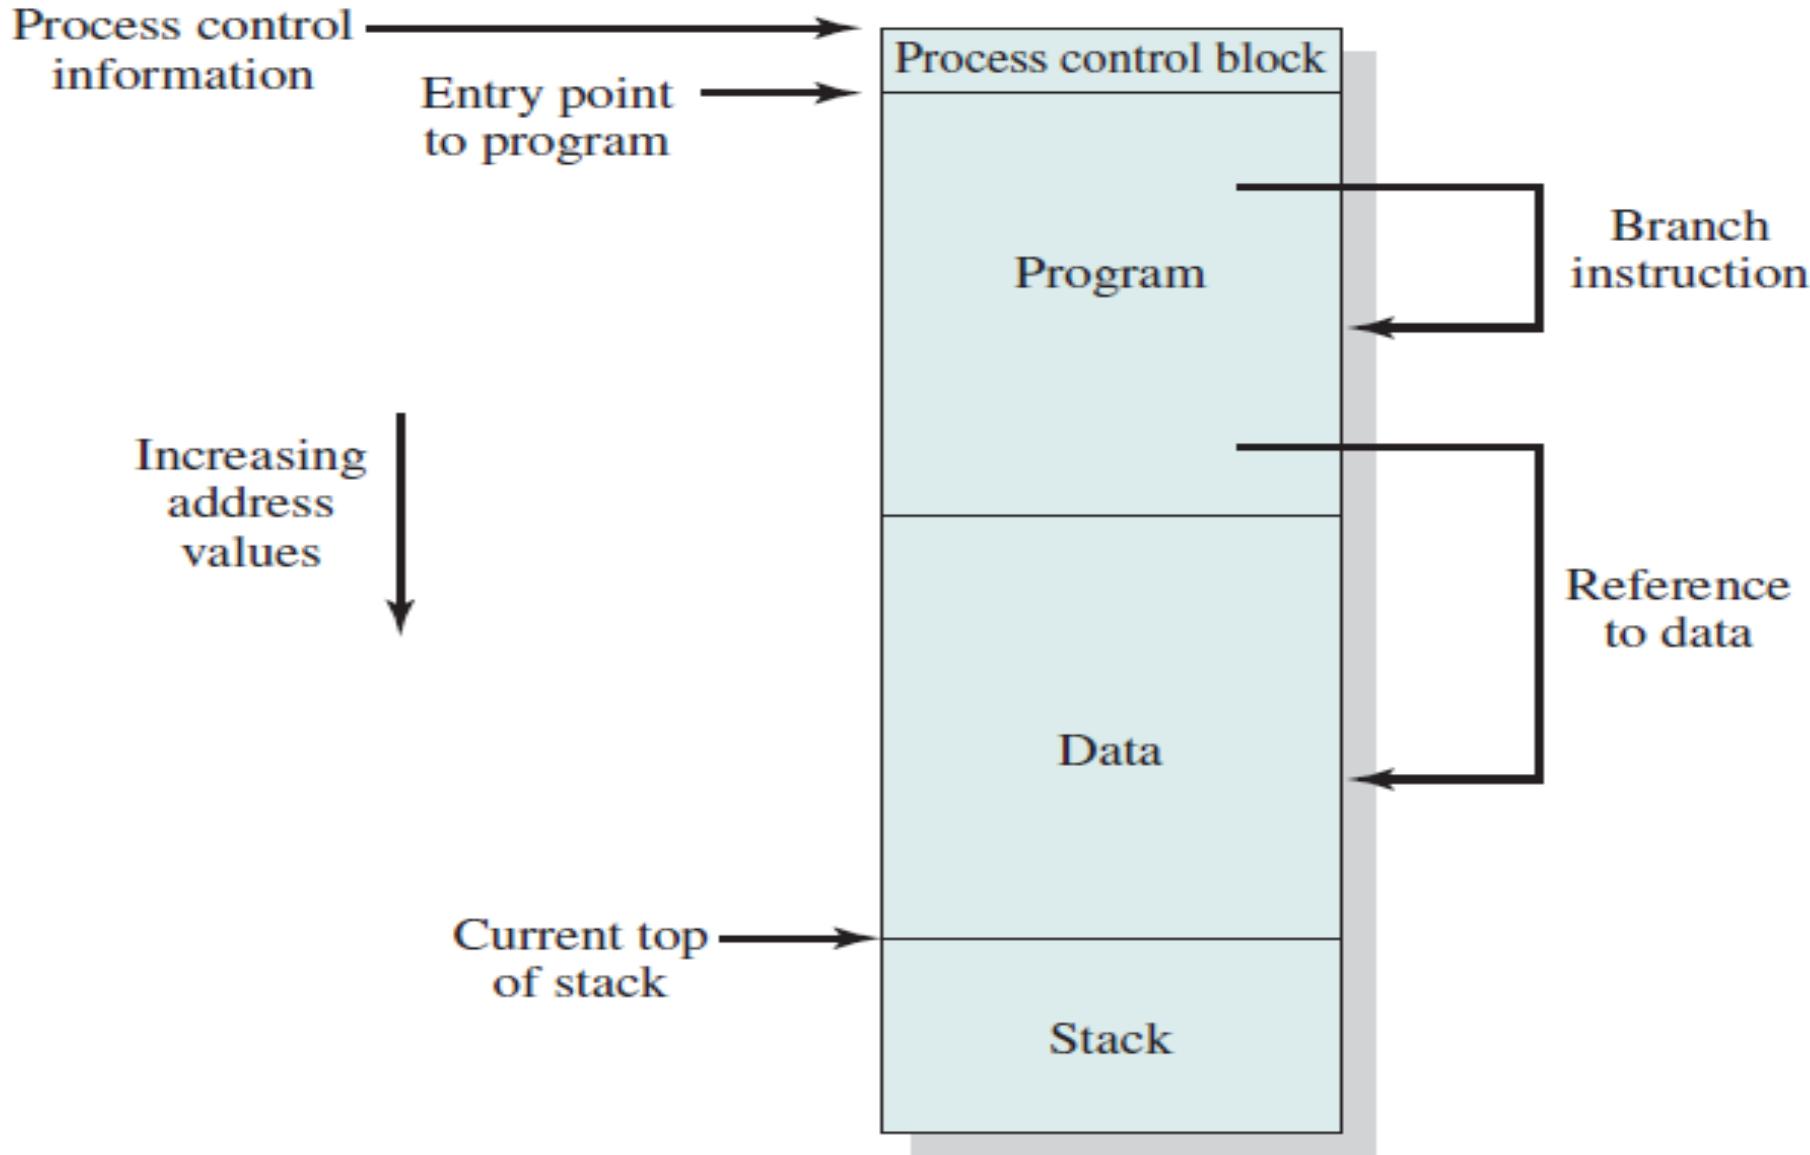
## OS provides API functions to

- Allow program to allocate, access and return memory

## When memory is allocated to processes, OS must keep track of

- which parts of the memory are free
- Which parts of the memory are allocated
- Allocated to which processes

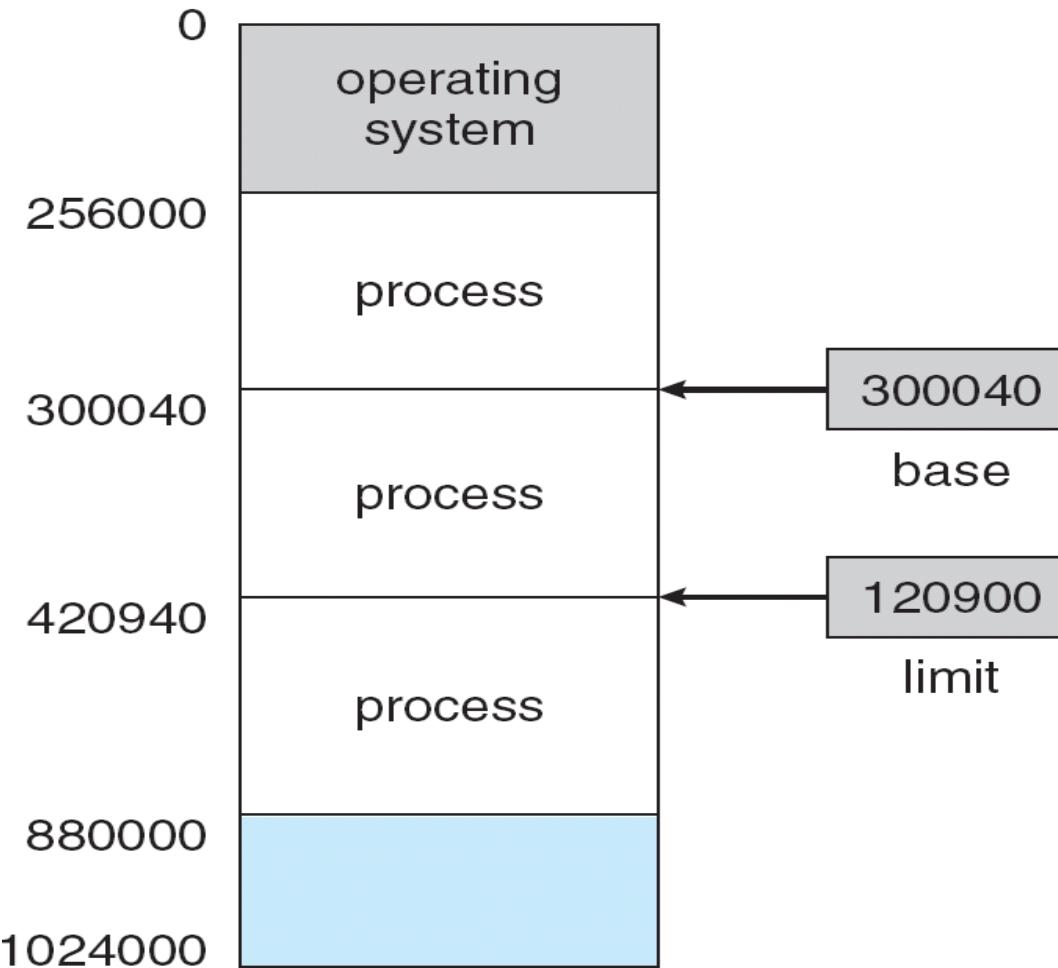
# Process in Memory



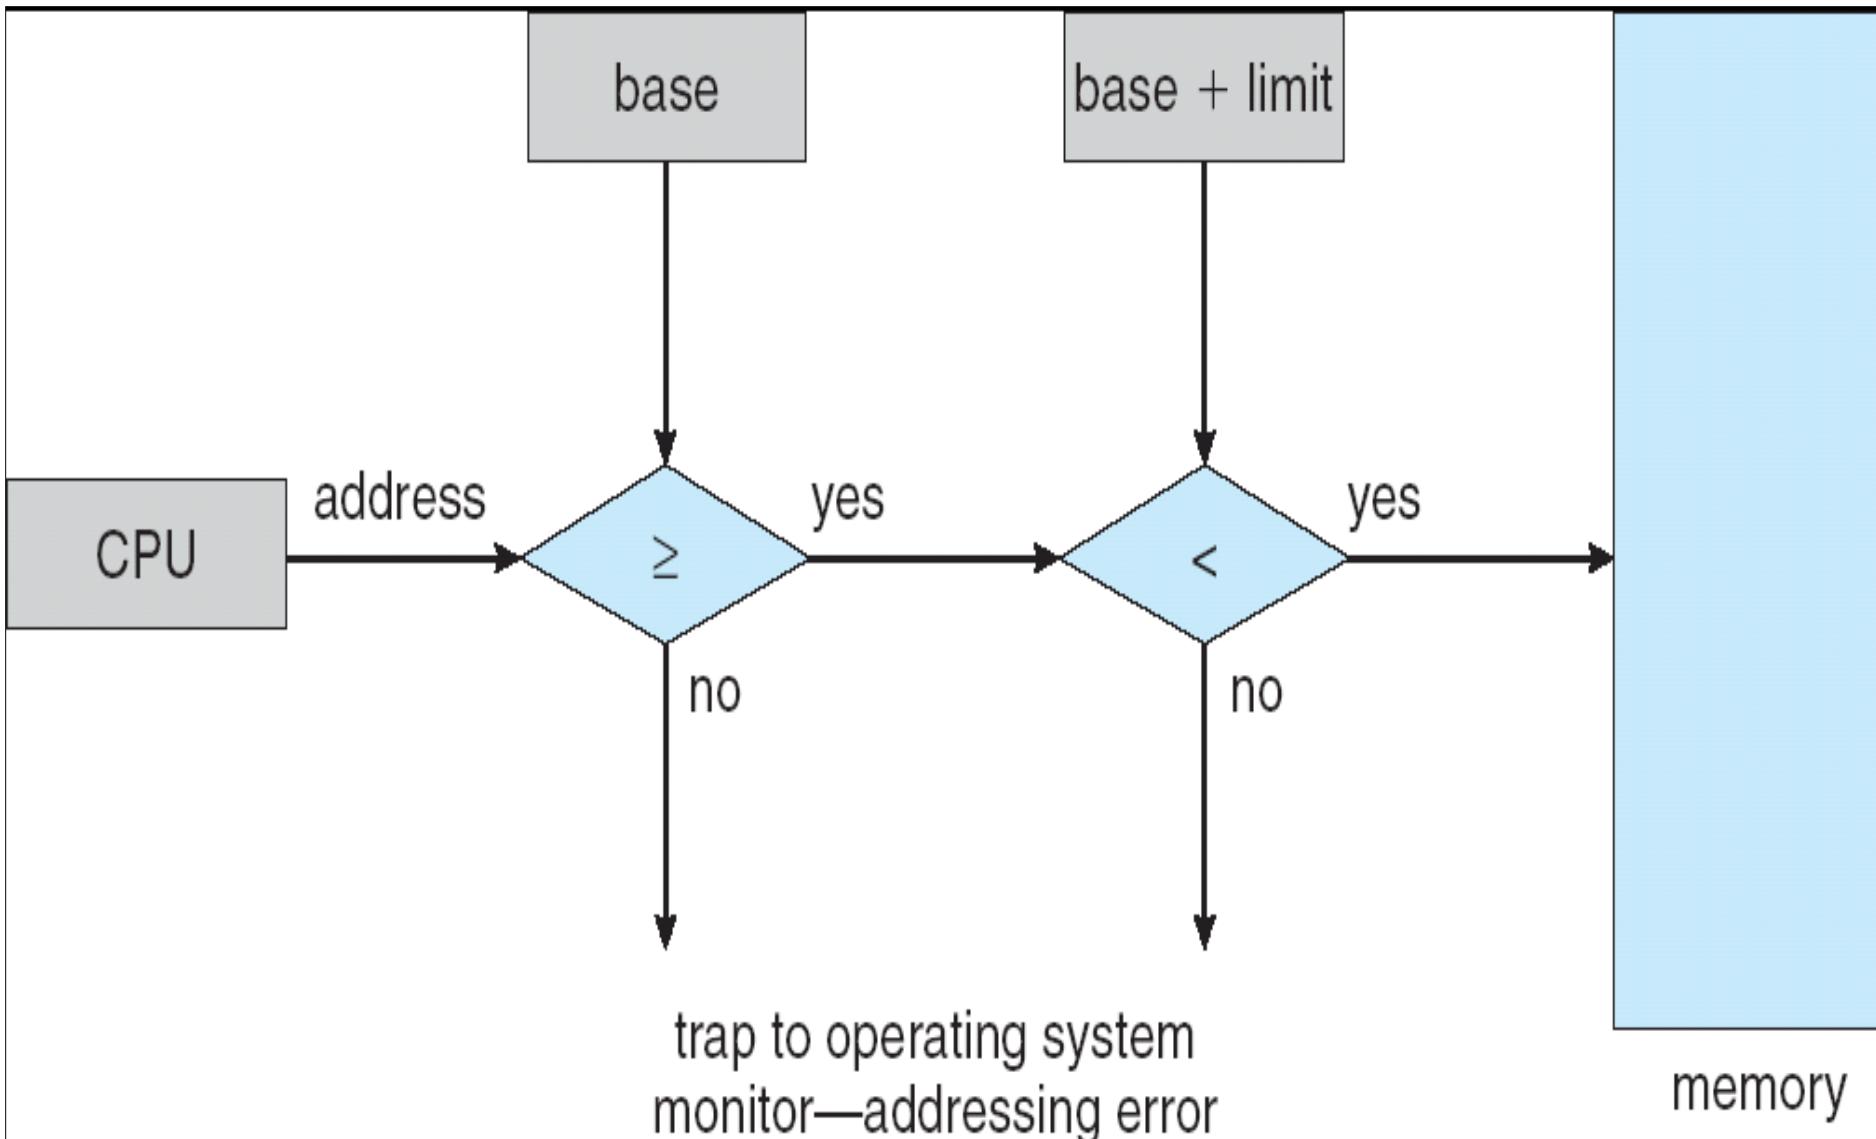
# Base and Limit Registers

---

- A pair of **base** and **limit** registers define the logical address space

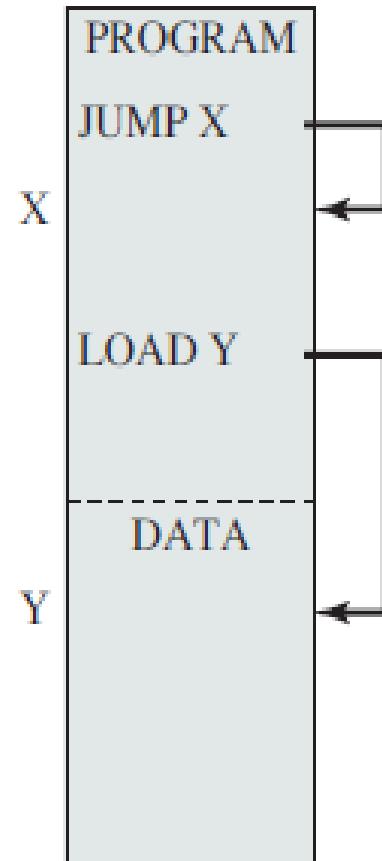


# HW address protection with base and limit registers

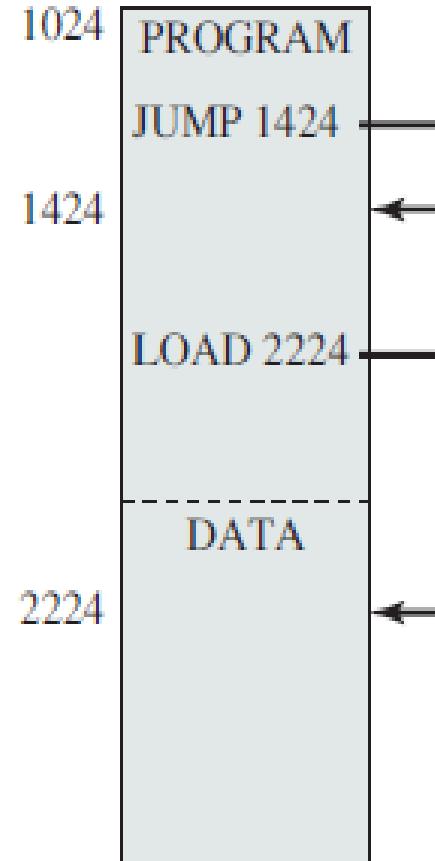


# Absolute and relative address

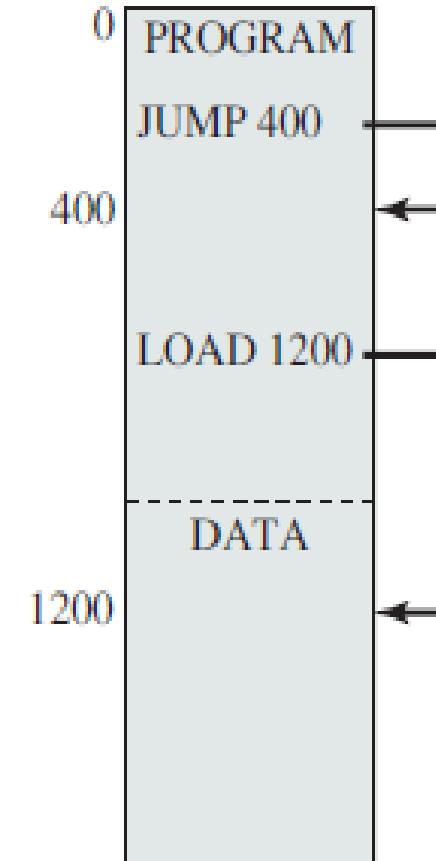
Symbolic  
addresses



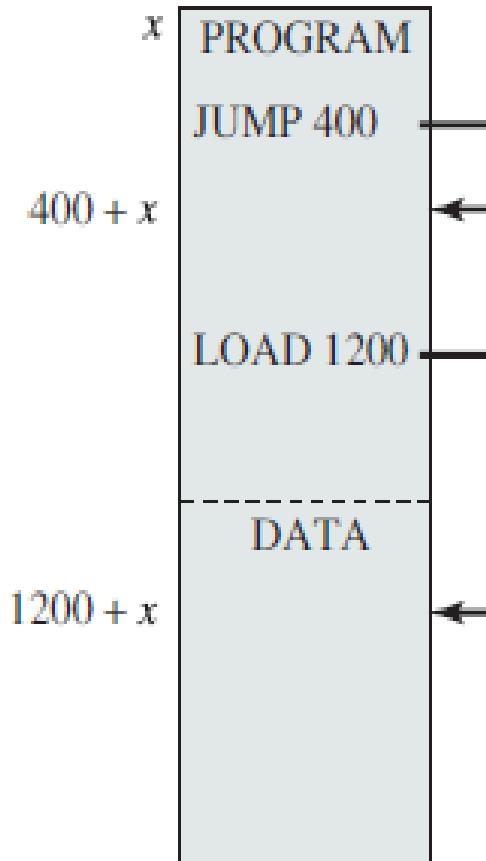
Absolute  
addresses



Relative  
addresses



Main memory  
addresses



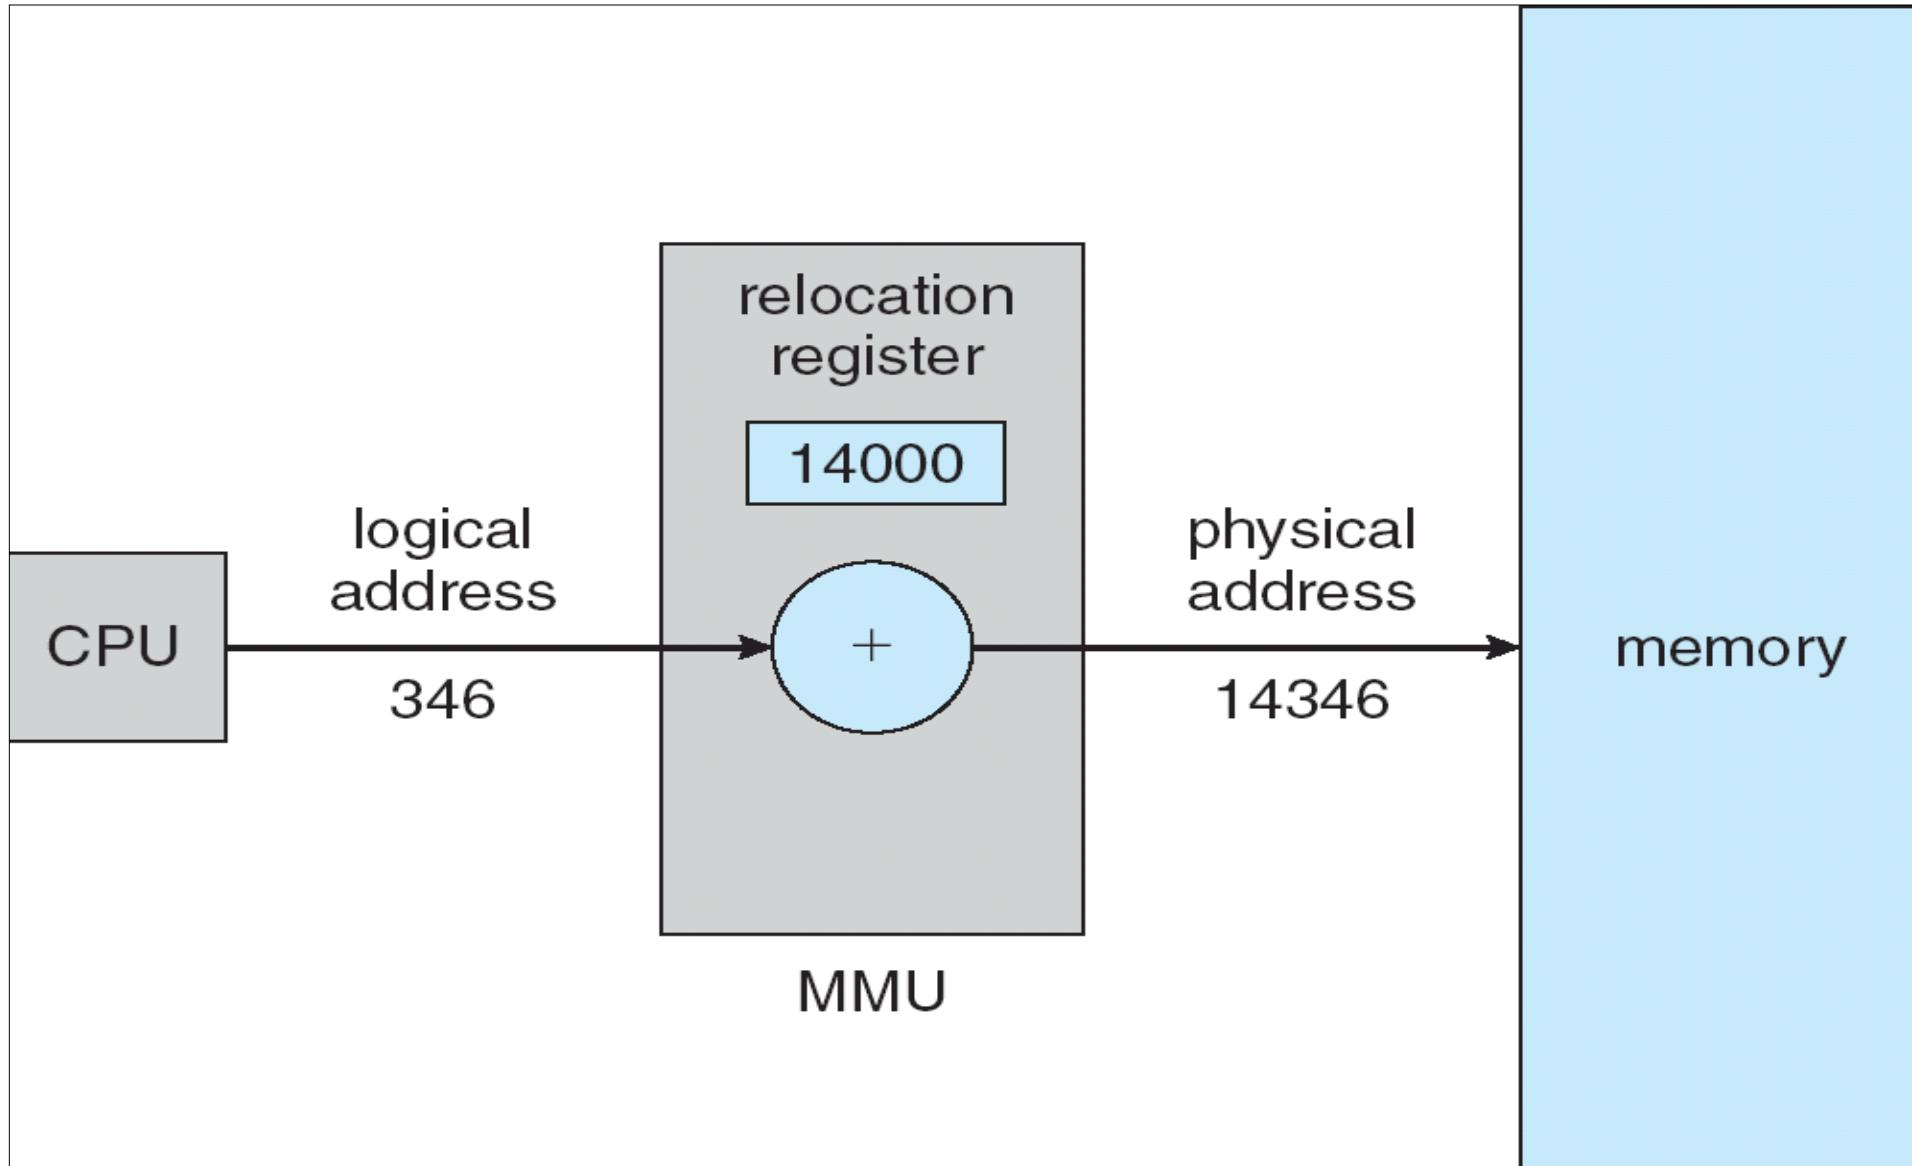
(a) Object module

(b) Absolute load module

(c) Relative load module

(d) Relative load module  
loaded into main memory  
starting at location x

# Dynamic relocation using a relocation register



# Memory Management Requirements

## Relocation

- To support multiprogramming

## Protection

- Against unwanted interferences

## Sharing

- Flexibility

## Logical organization

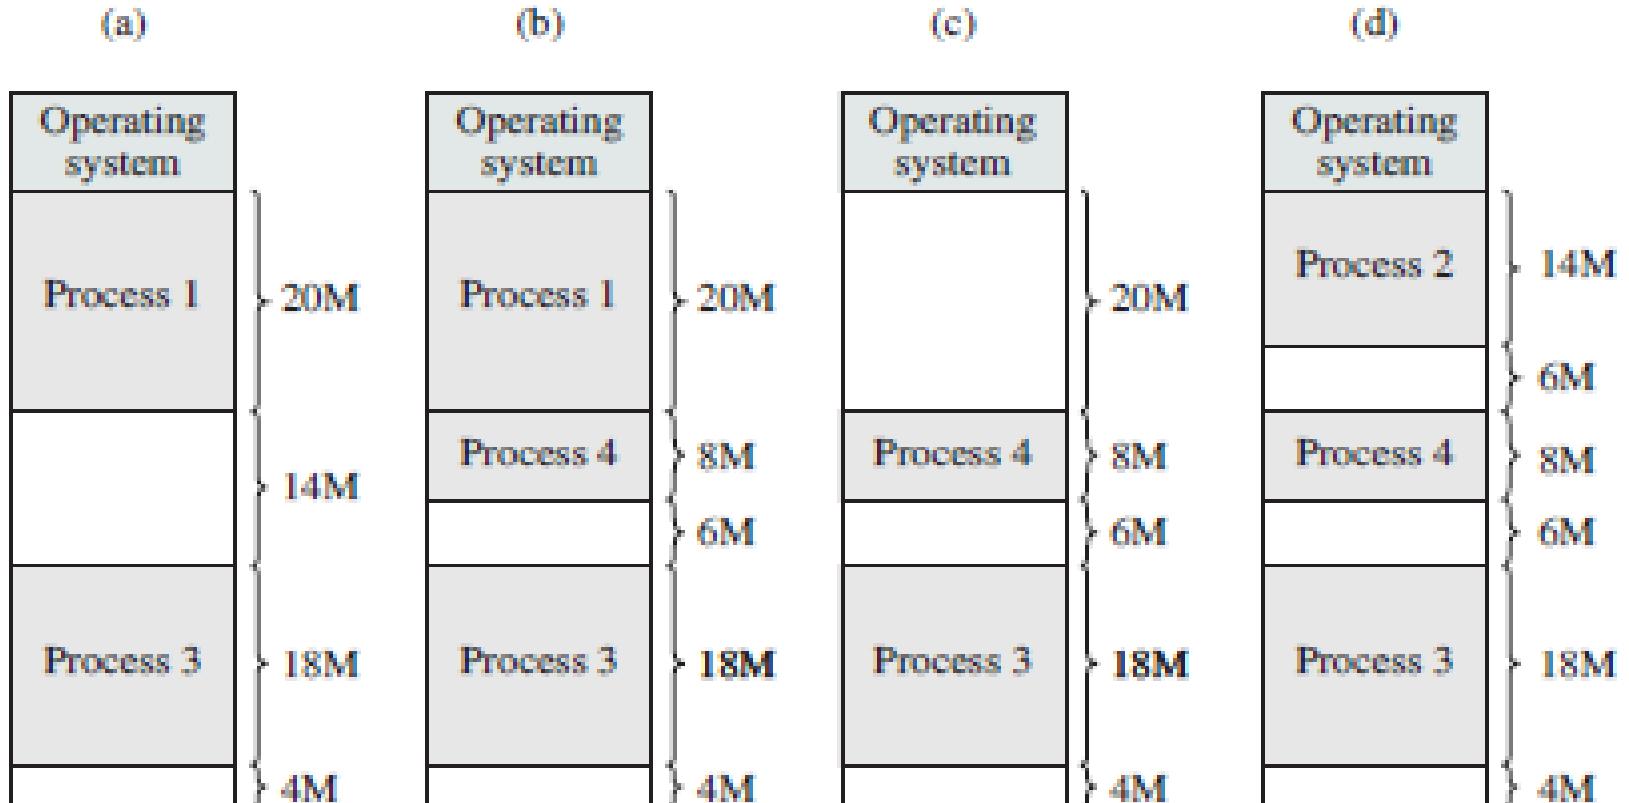
- Linear address space

## Physical organization

- Space is limited

# Contiguous Memory Allocation – dynamic partition

- Creating partitions dynamically of exactly the same size as the process
- No internal fragmentation
- Problem: External fragmentation



# Fragmentation

---

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time

# Paging

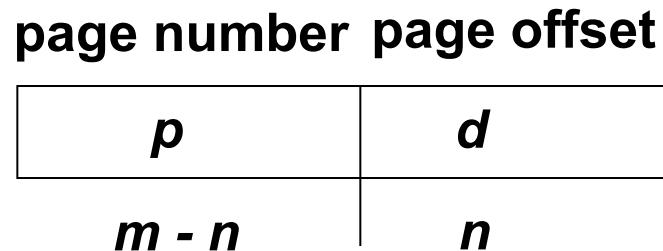
---

- Noncontiguous memory allocation
- Divide physical memory into fixed-sized blocks called **frames**
- Divide logical memory into blocks of same size called **pages**
- Set up a page table to translate logical to physical addresses
- To run a program of size  $n$  pages, need to find  $n$  free frames and load program
- Must keep track of all free frames
- **Problem:** internal fragmentation

# Address Translation Scheme

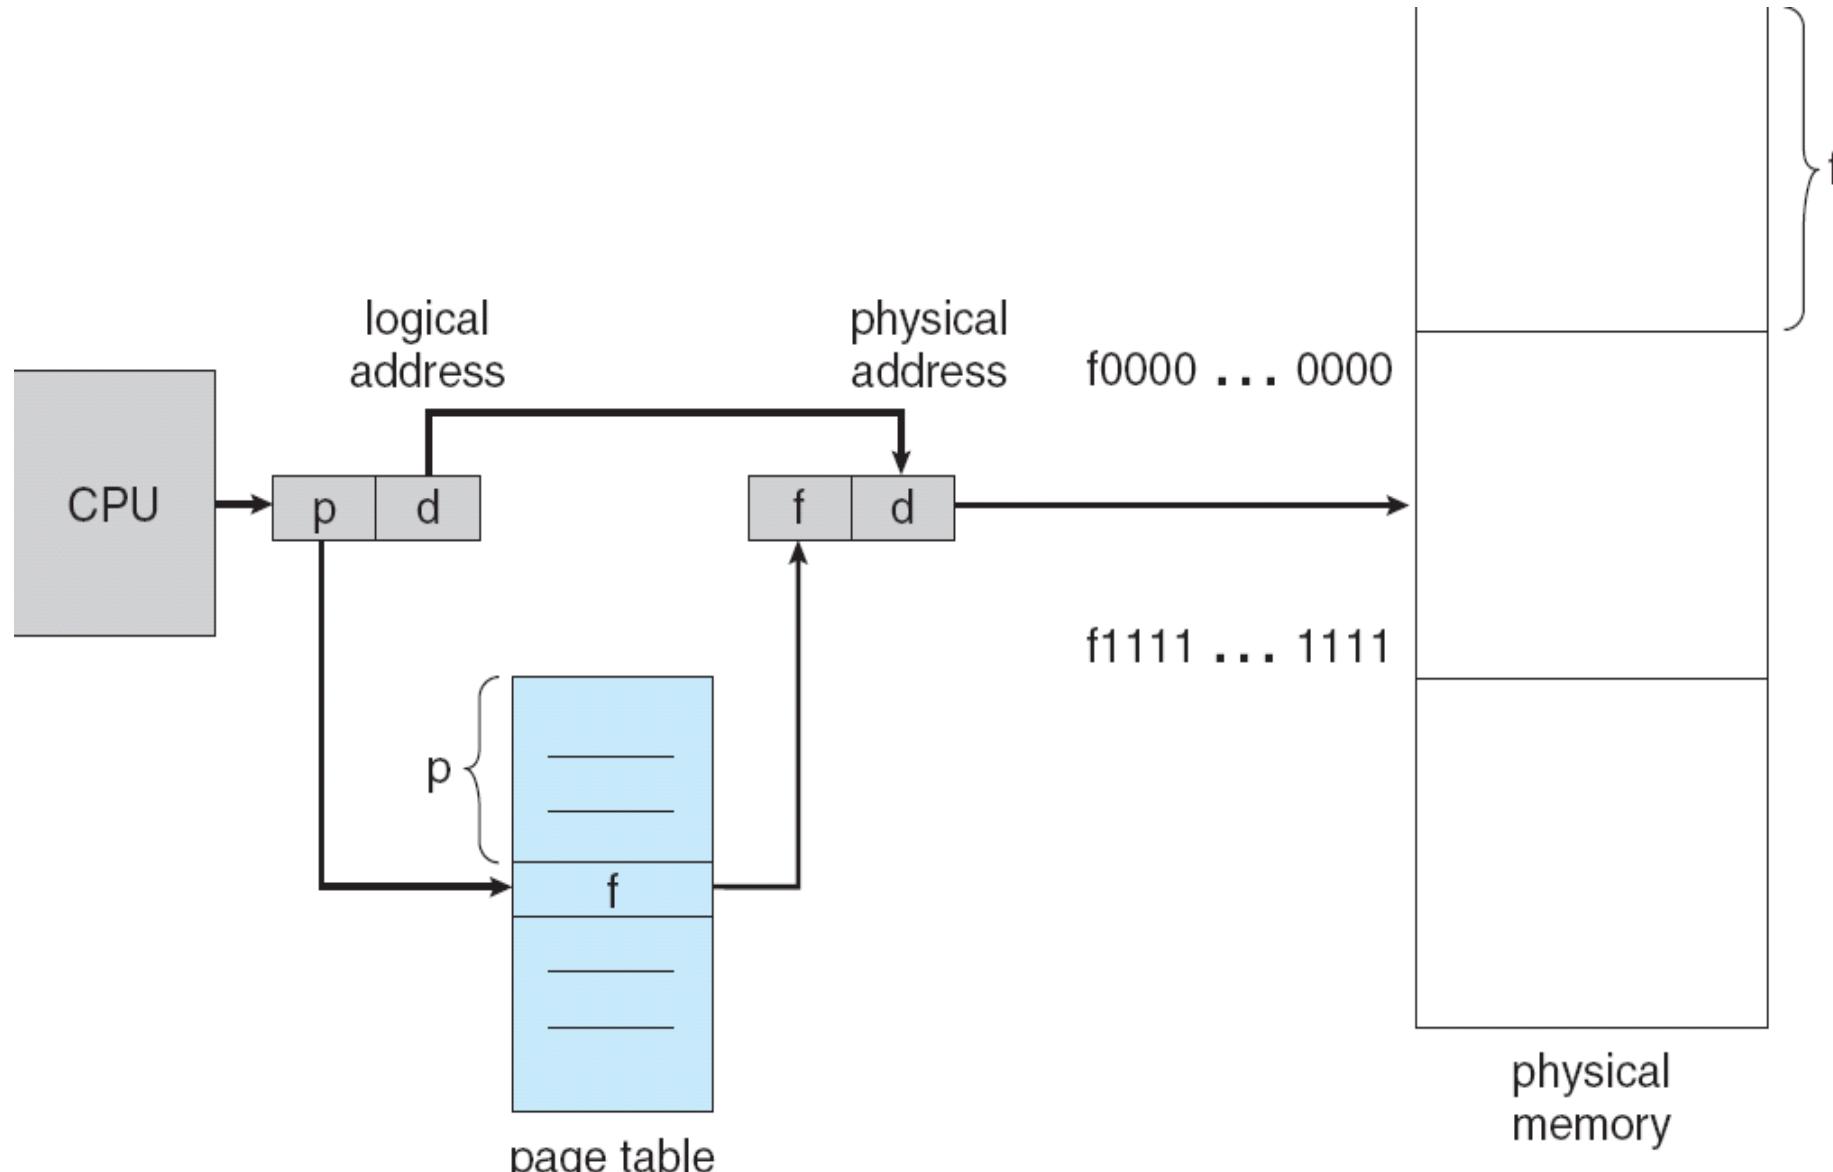
---

- Address generated by CPU is divided into:
  - **Page number ( $p$ )** – used as an index into a *page table* which contains base address of each page in physical memory
  - **Page offset ( $d$ )** – combined with base address to define the physical memory address that is sent to the memory unit

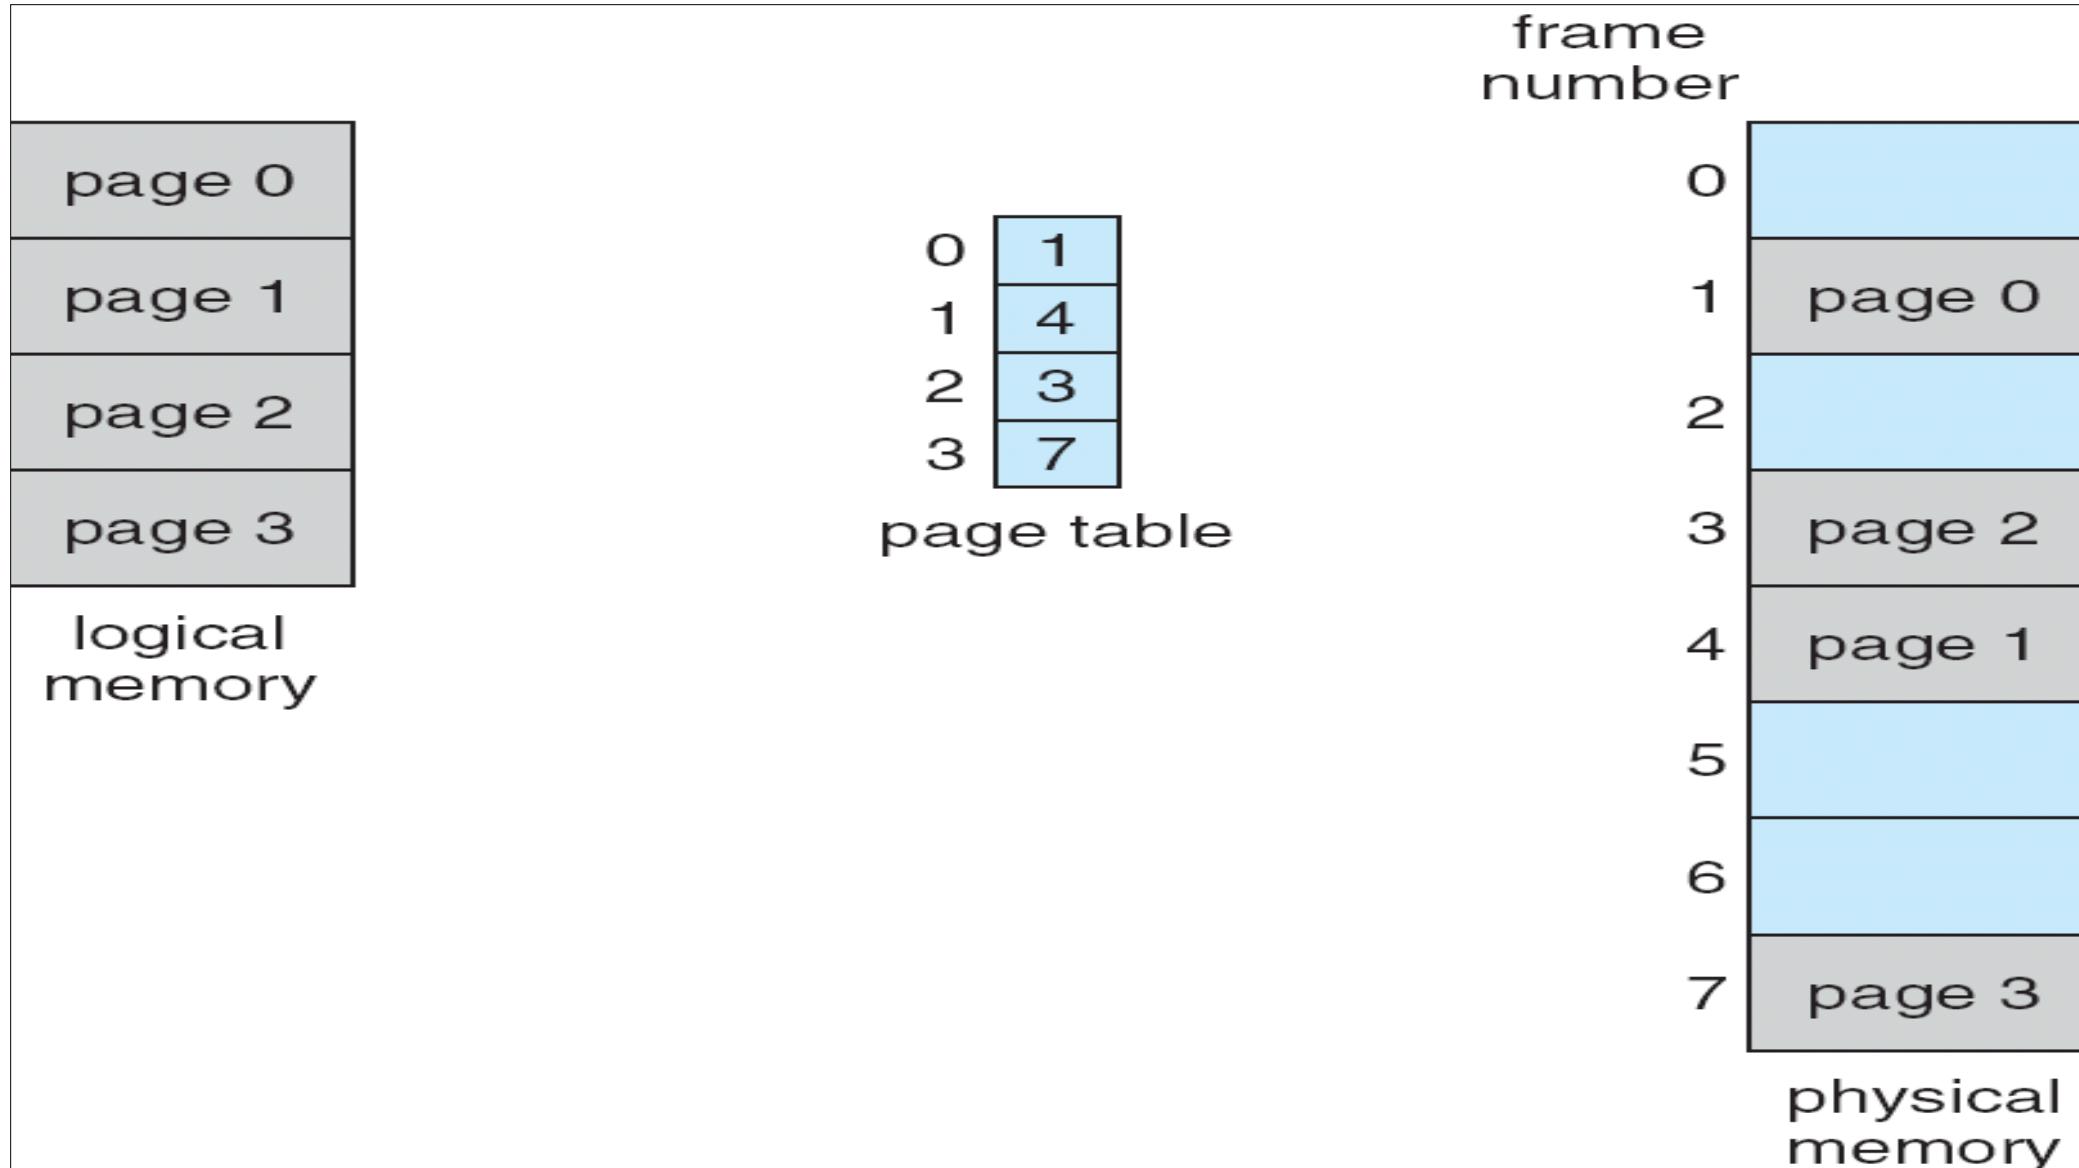


- For given logical address space  $2^m$  and page size  $2^n$

# Paging Hardware



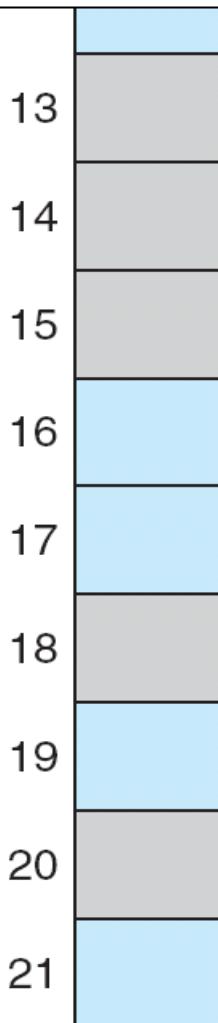
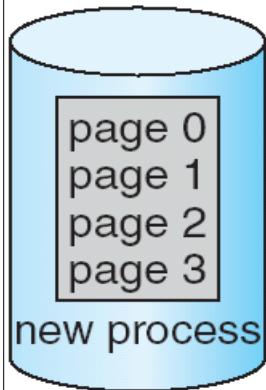
# Paging Model of Logical and Physical Memory



# Free Frames

free-frame list

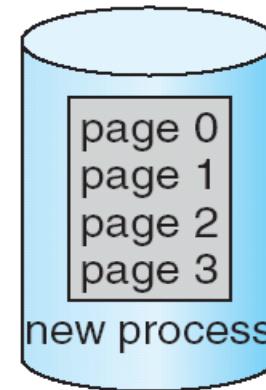
14  
13  
18  
20  
15



(a)

free-frame list

15



new-process page table

0	14
1	13
2	18
3	20

(b)

Before allocation

After allocation

# Implementation of Page Table

---

- Scheme #1
  - Page table is kept in main memory
  - **Page-table base register (PTBR)** points to the page table
  - **Page-table length register (PRLR)** indicates size of the page table
  - In this scheme every data/instruction access requires two memory accesses. One for the page table and one for the data/instruction.
- Scheme #2
  - The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory or translation look-aside buffers (TLBs)**

# Translation look-aside buffers (TLBs)

---

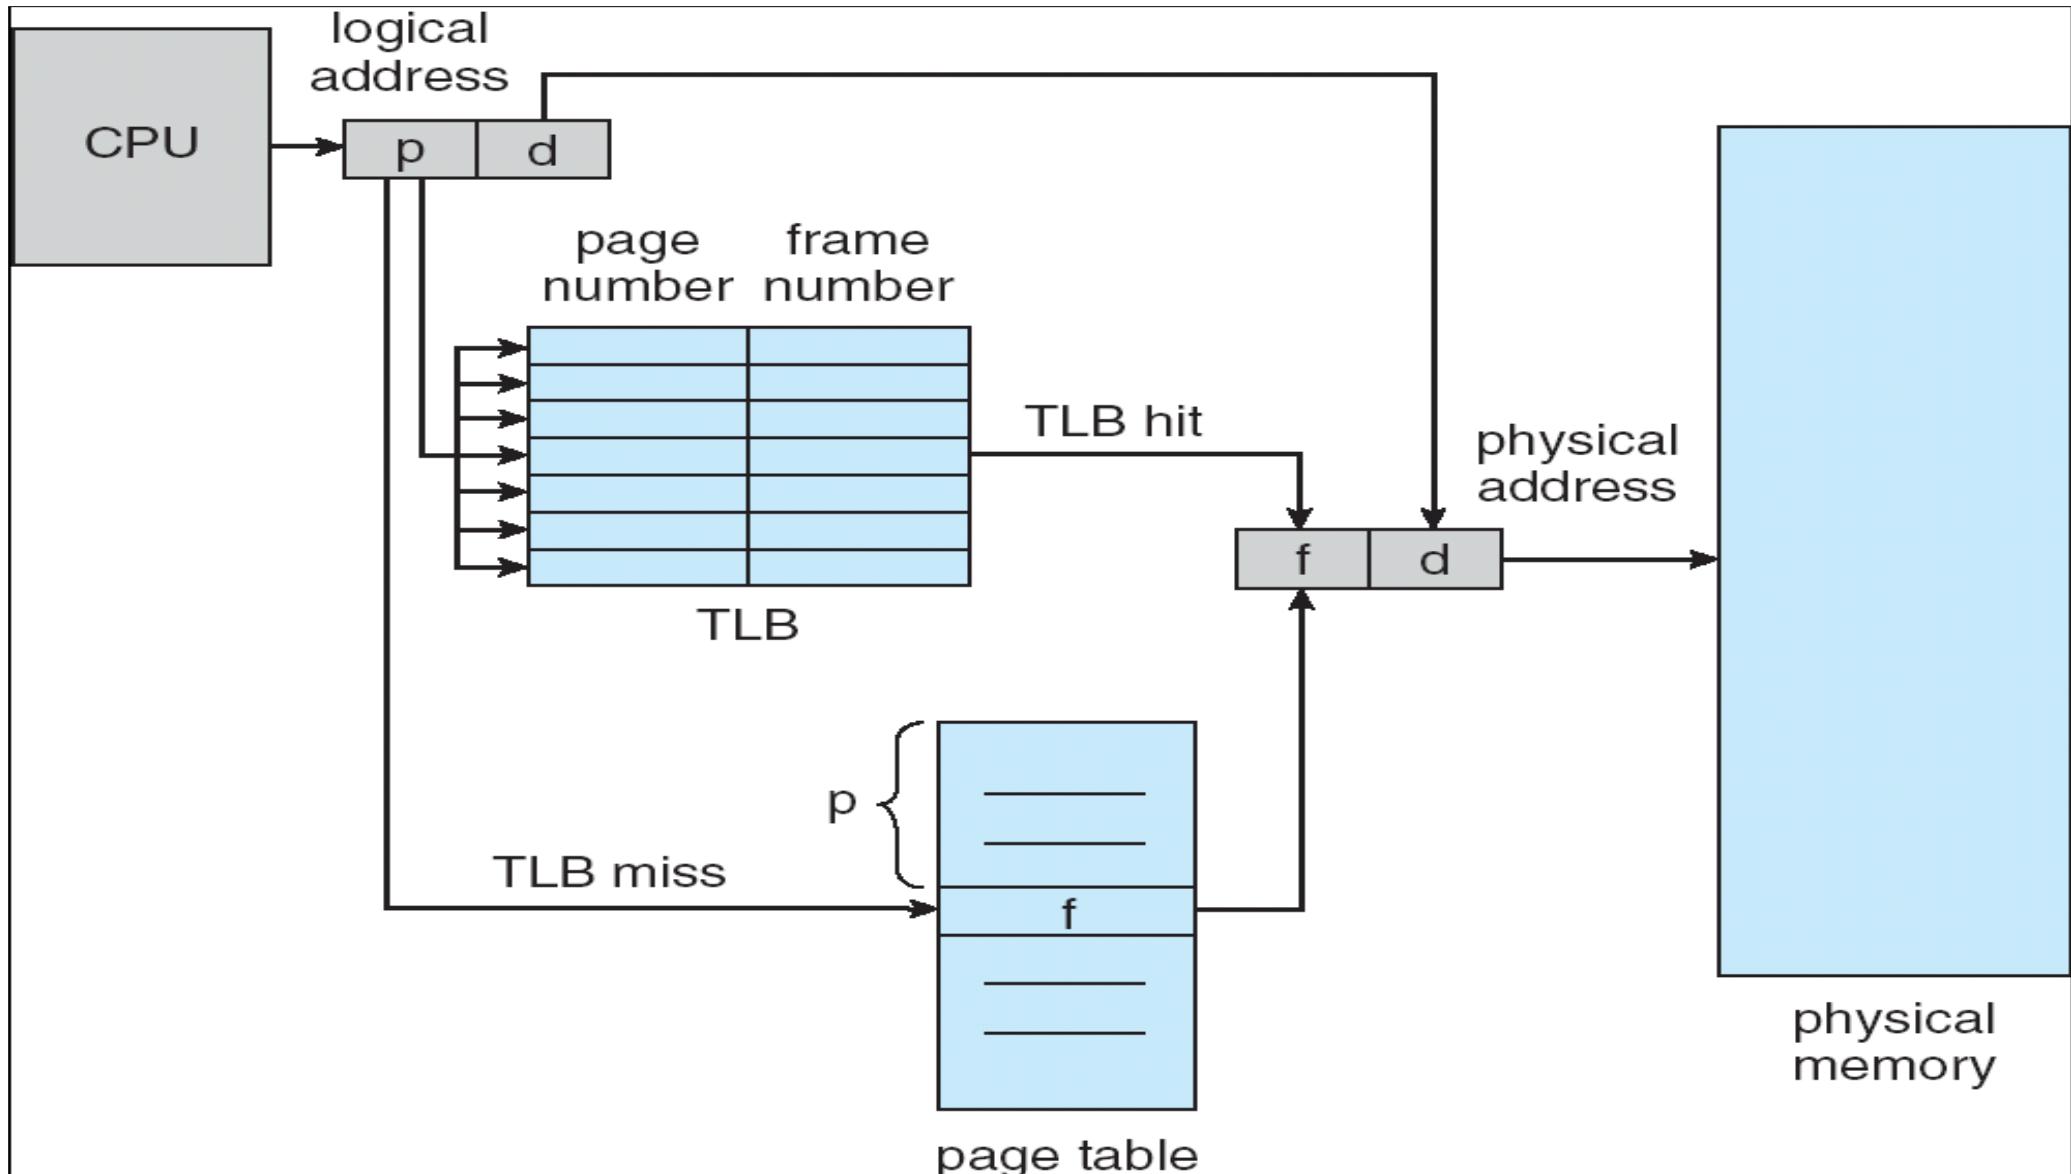
- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
  - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access

Page #	Frame #

- Associative memory – parallel search
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory

Address translation (p, d)

# Paging Hardware With TLB



# Effective Access Time

---

## □ **Effective Access Time (EAT)**

$$EAT = (TLB_h * HitR) + (TLB_m * (1 - HitR))$$

- Hit ratio (HitR) – percentage of times that a page number is found in the associative registers
- Time to access TLB and access memory when TLB hit (TLB<sub>h</sub>)
- Time to access TLB and access page table in memory and access memory when TLB miss (TLB<sub>m</sub>)

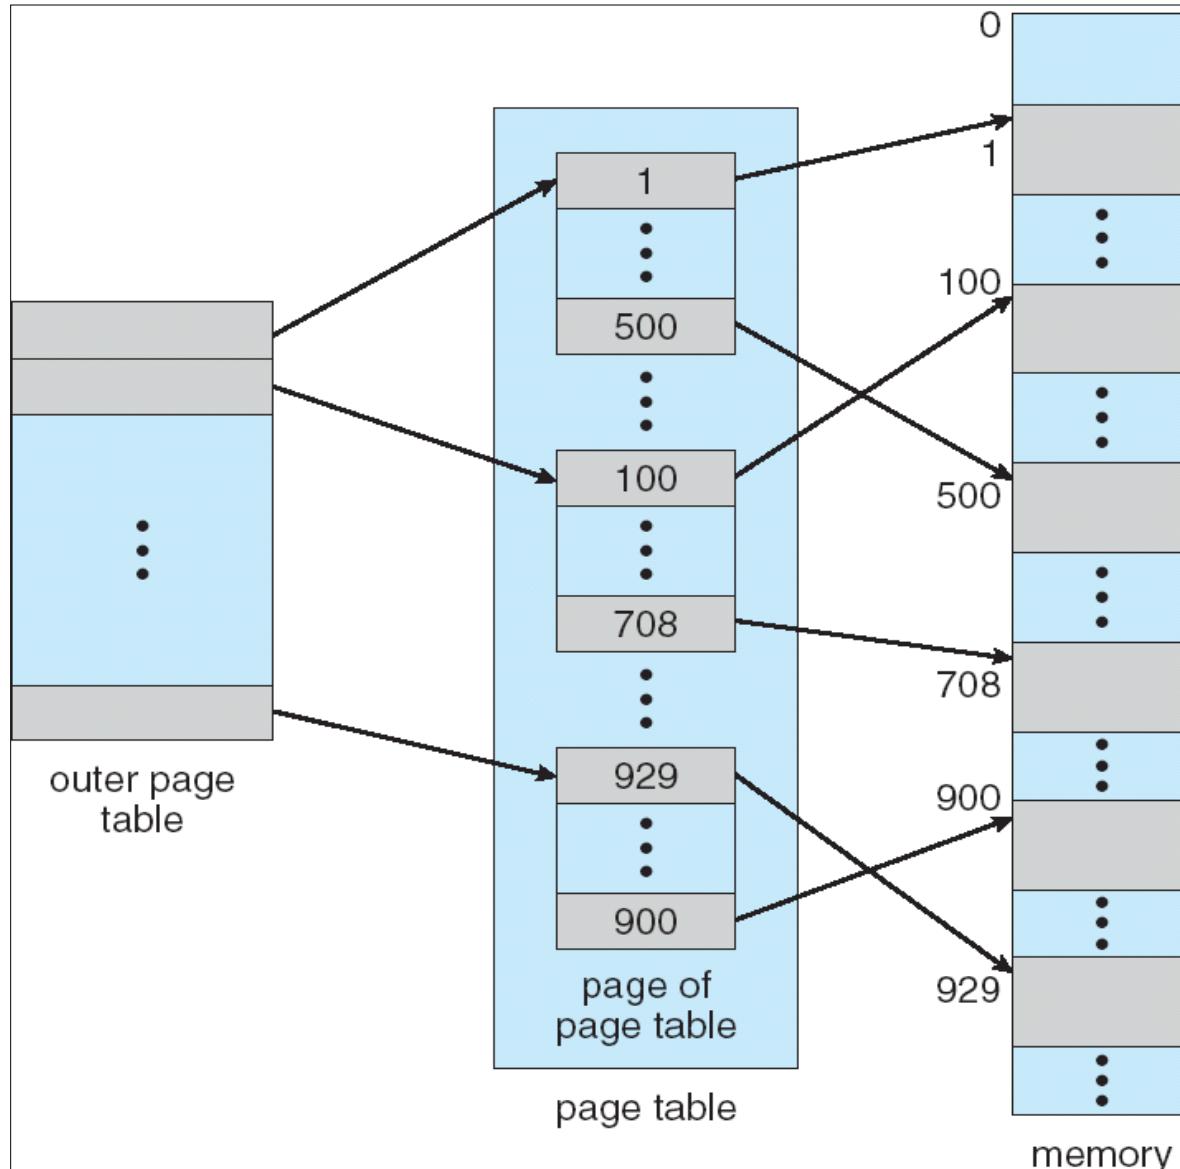
# Effective Access Time (example)

---

- Consider HitR = 80%,  $\varepsilon$  = 20ns for TLB search, 100ns for memory access
  - EAT =  $0.80 \times 120 + 0.20 \times 220$   
= 96 + 44  
= 140ns
  
- Consider more realistic hit ratio -> HitR = 99%,  $\varepsilon$  = 20ns for TLB search, 100ns for memory access
  - EAT =  $0.99 \times 120 + 0.01 \times 220$   
= 118.8 + 2.2  
= 121ns

$\varepsilon$  = Time for TLB search

# Hierarchical Page Tables



- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table

# Two-Level Paging Example

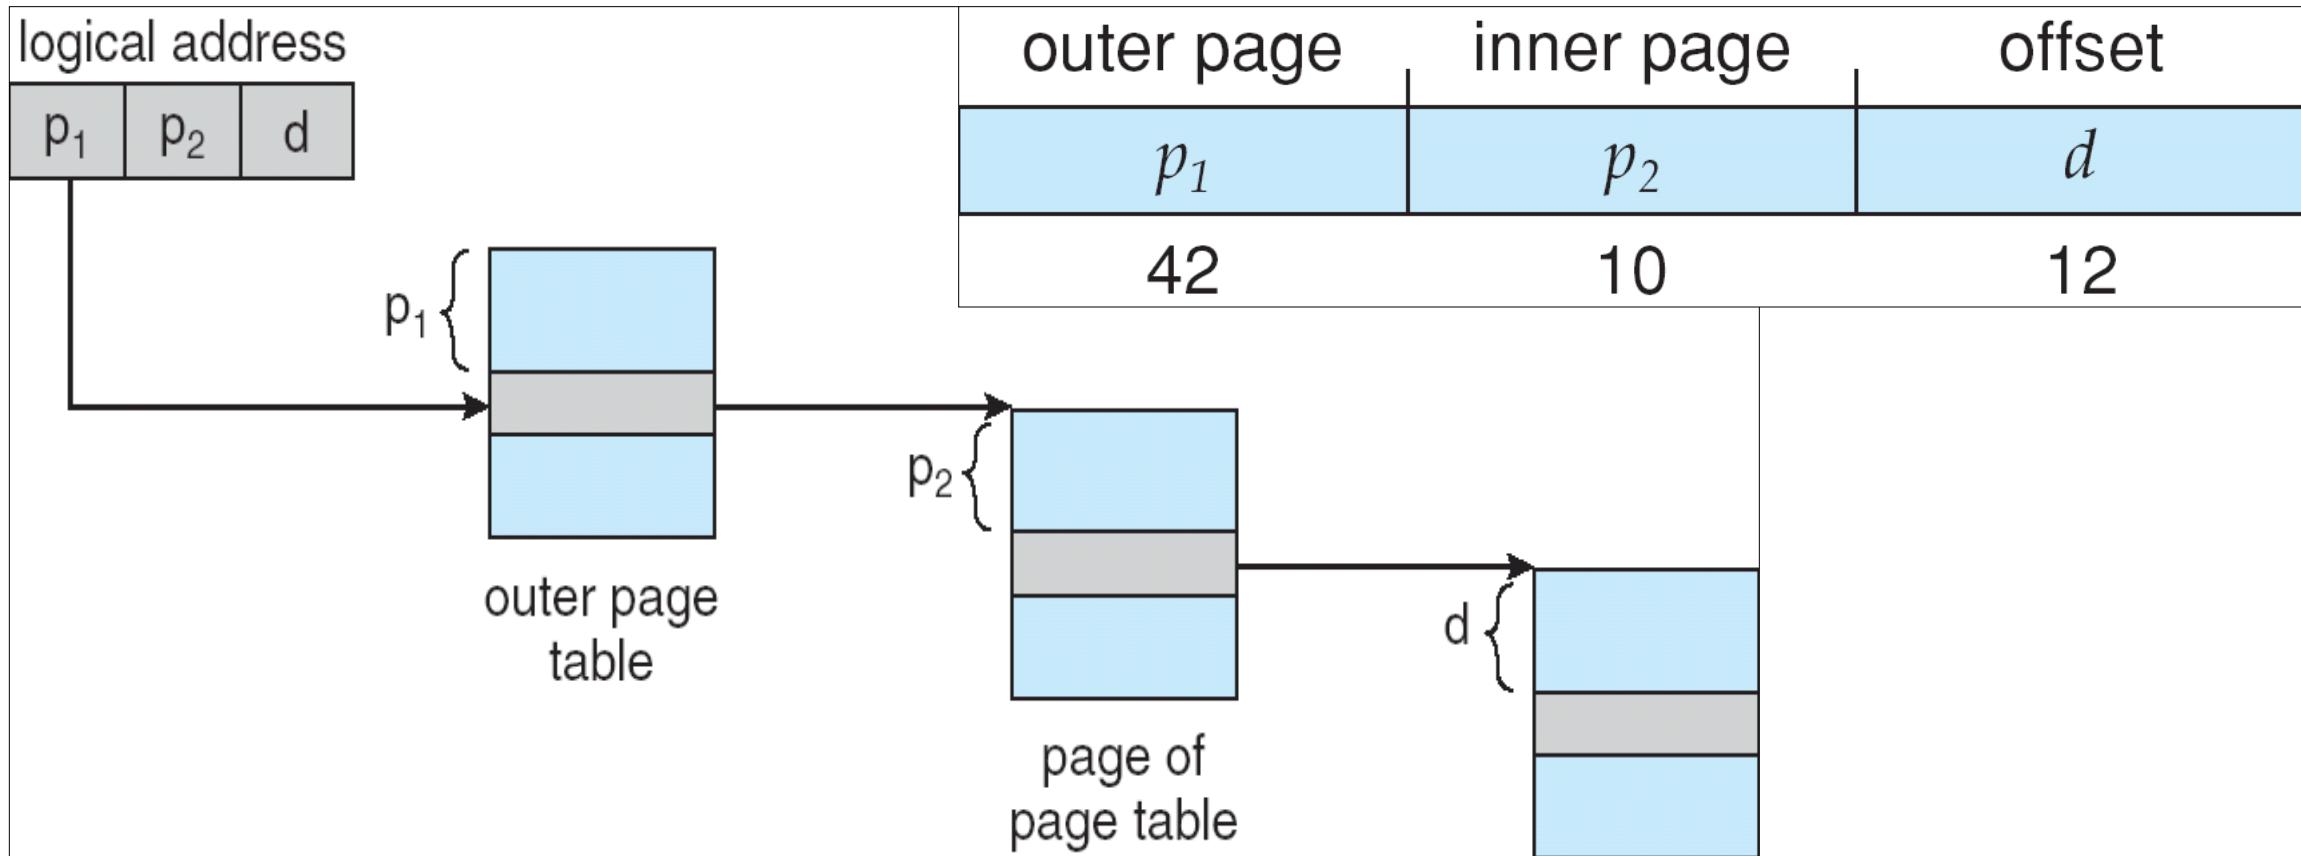
---

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:

page number	page offset
$p_i$	$p_2$
12	10

where  $p_i$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the outer page table

# Address-Translation Scheme



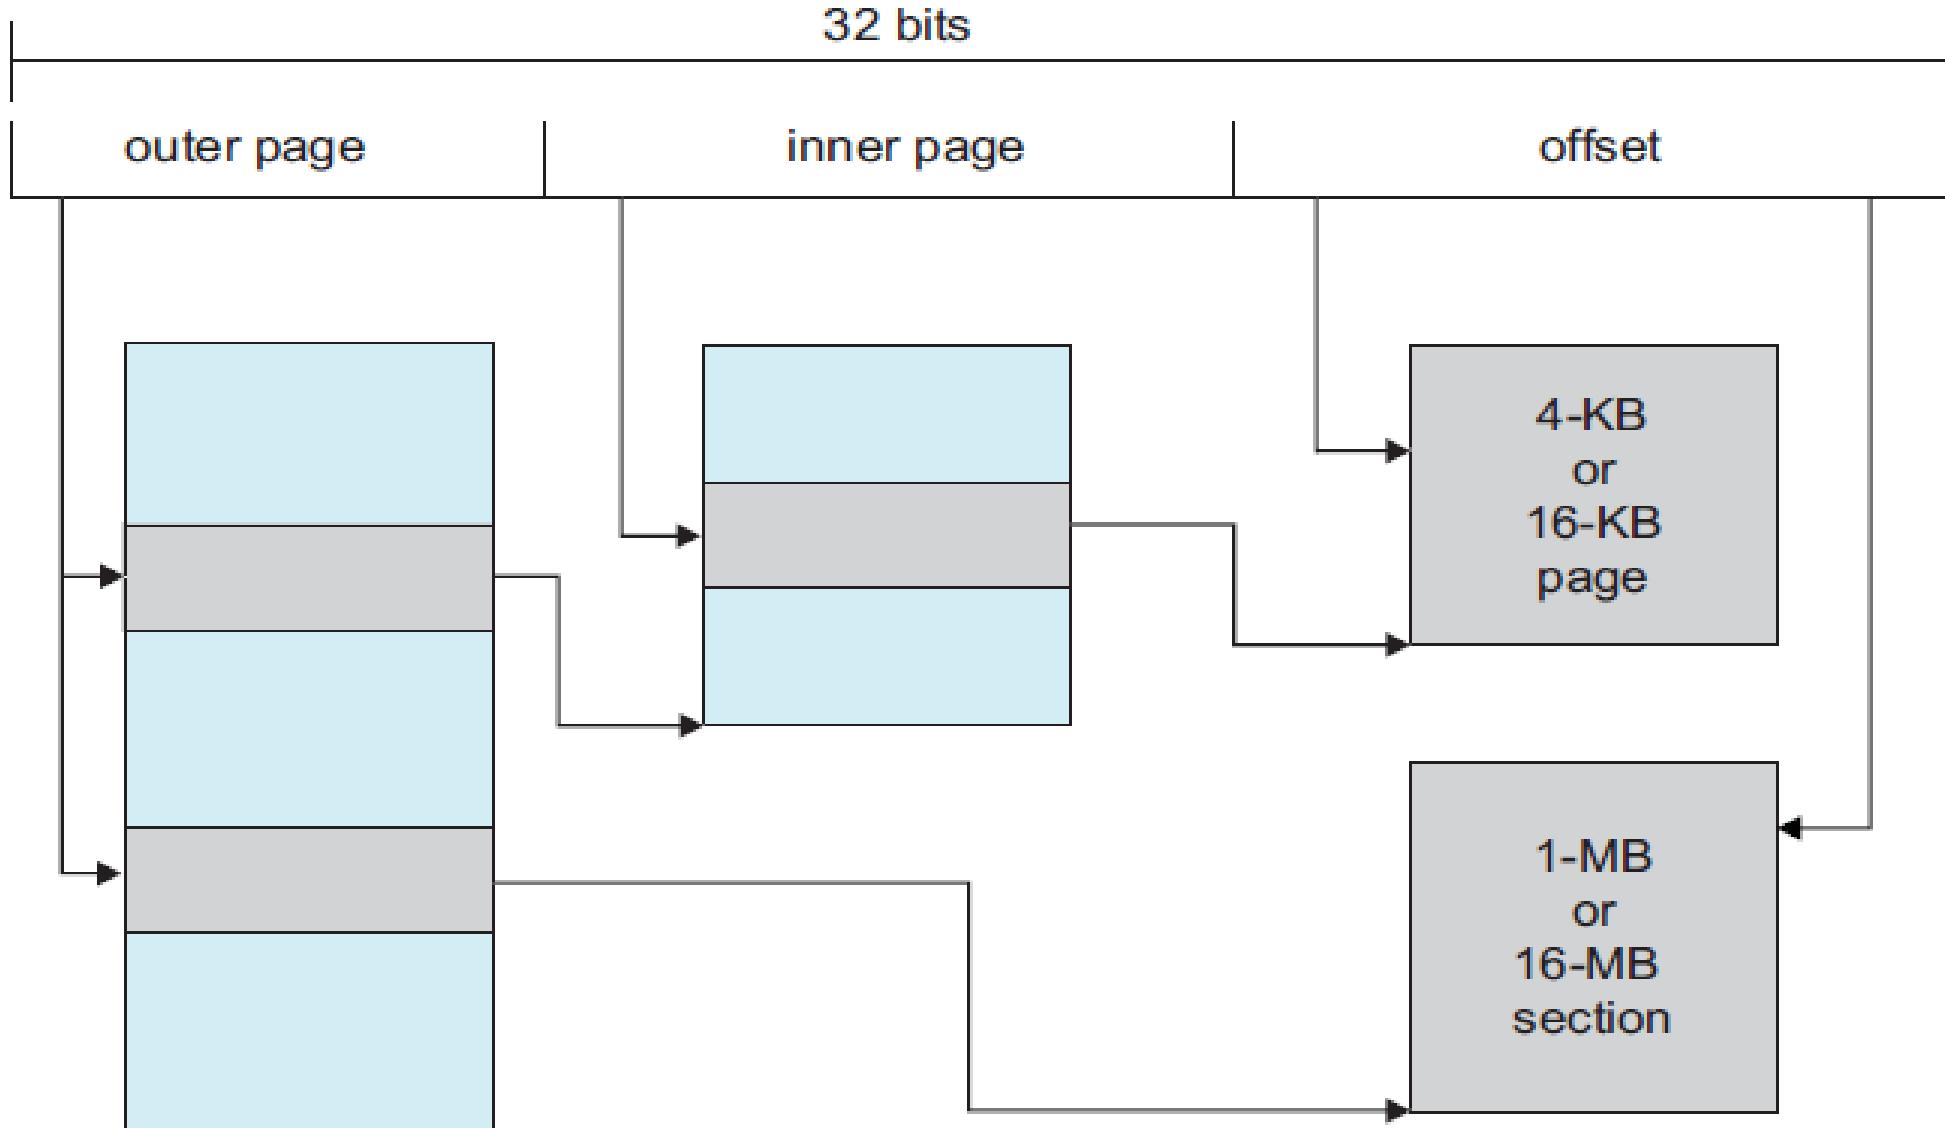
2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12

# Paging example: ARM

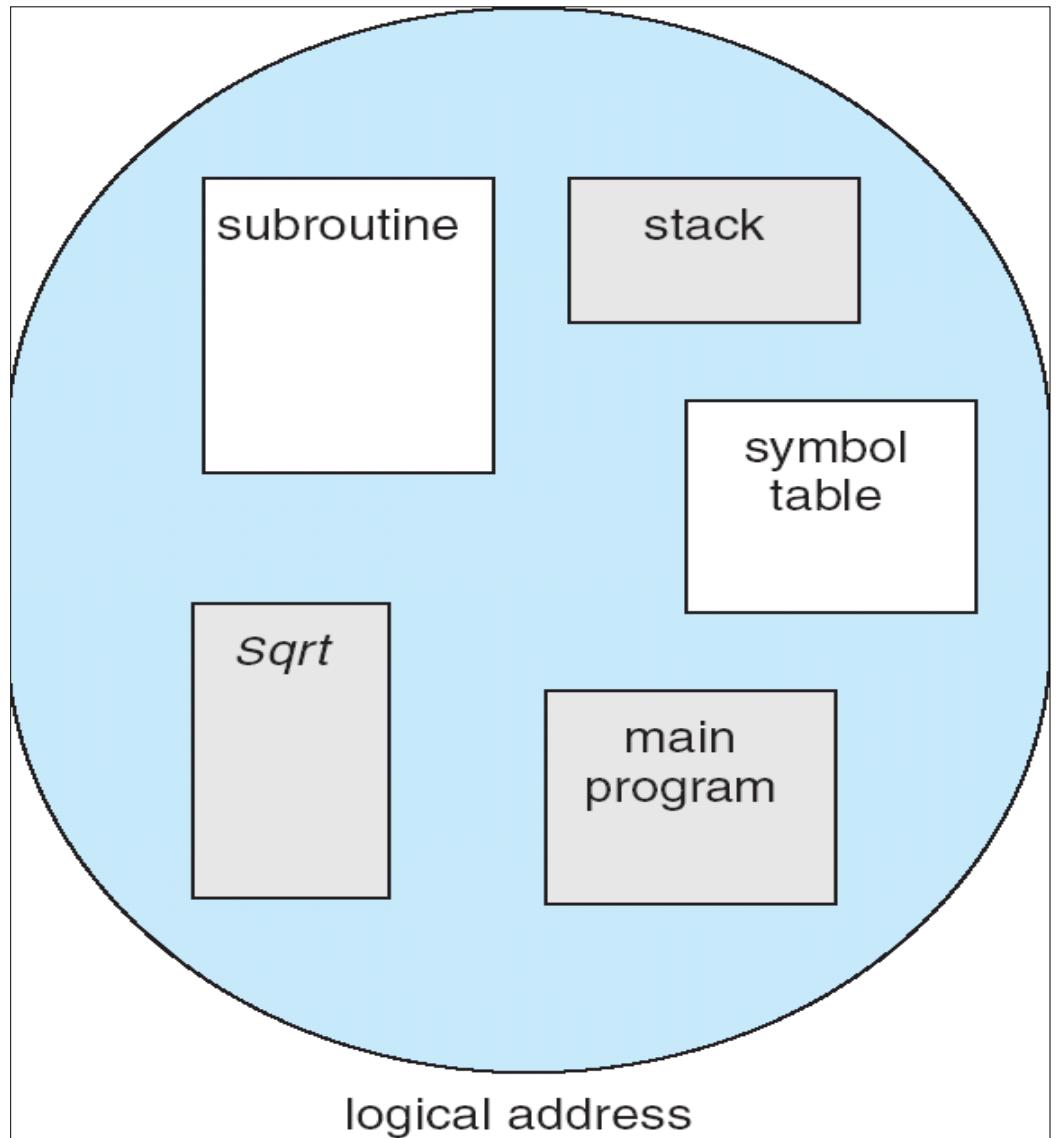
---

- Dominant mobile platform chip
  - Apple iOS and Google Android devices
- Modern, energy efficient, 32-bit CPU
  - 4 KB and 16 KB pages
  - 1 MB and 16 MB pages (termed **sections**)
- One-level paging for sections, two-level paging for smaller pages
- Two levels of TLBs
  - Outer level has two micro TLBs (data, instruction)
  - Inner is single main TLB
- First inner is checked, on missouters are checked, and on miss page table walk performed by CPU

# Paging example: ARM (cont.)



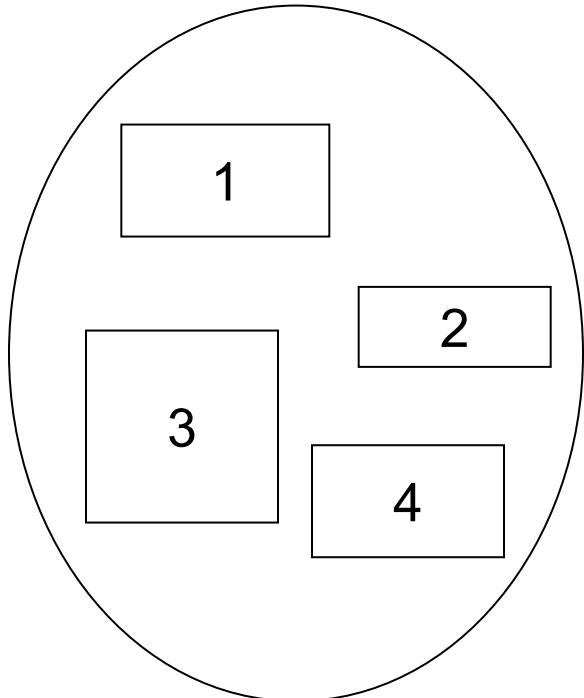
# Segmentation – user view



- Memory-management scheme that supports user view of memory
- A program is a collection of segments. A segment is a logical unit such as:
  - main program,
  - procedure,
  - function,
  - method,
  - object,
  - local variables, global variables,
  - common block,
  - stack,
  - symbol table, arrays

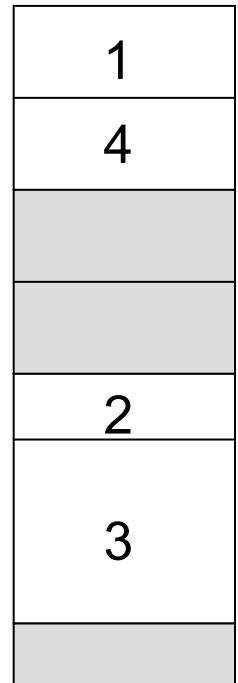
# Segmentation – logical view

- Memory-management scheme that supports user view of memory
- A program is a collection of segments. A segment is a logical unit such as:
  - main program,
  - procedure,
  - function,
  - method,
  - object,
  - local variables, global variables,
  - common block,
  - stack,
  - symbol table, arrays



user space

physical memory space

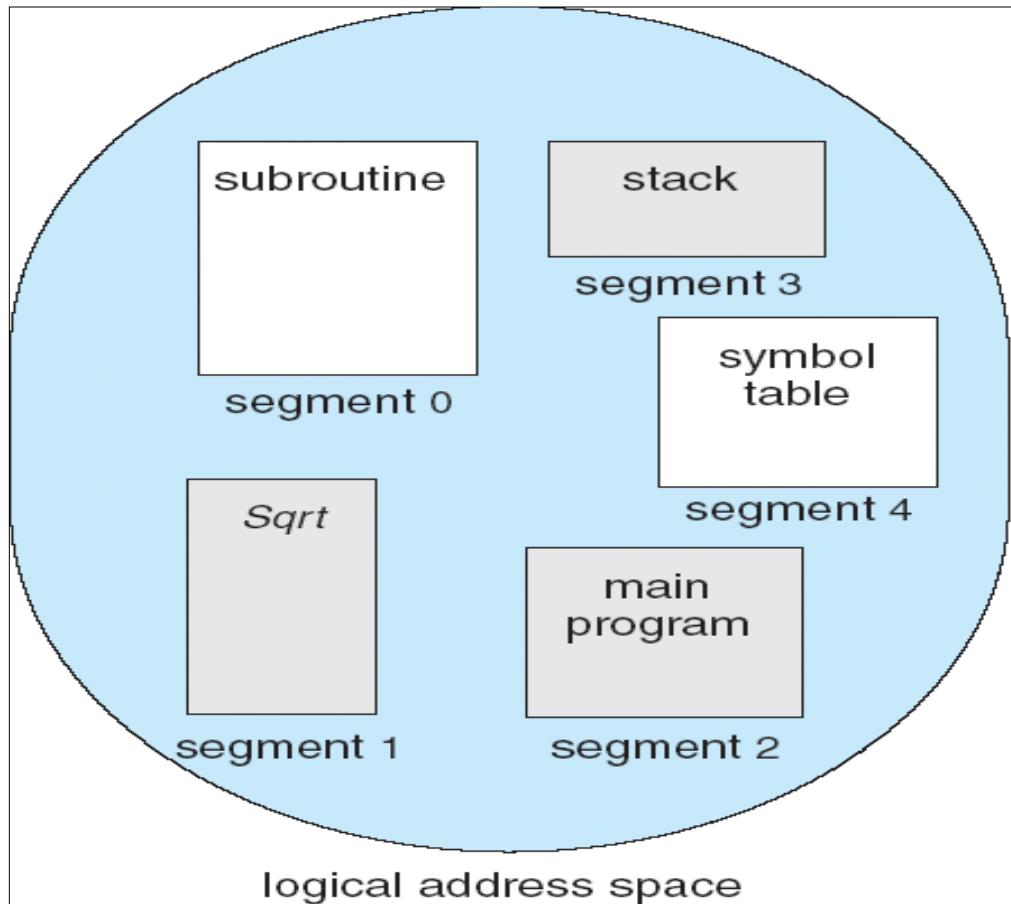


# Segmentation Architecture

---

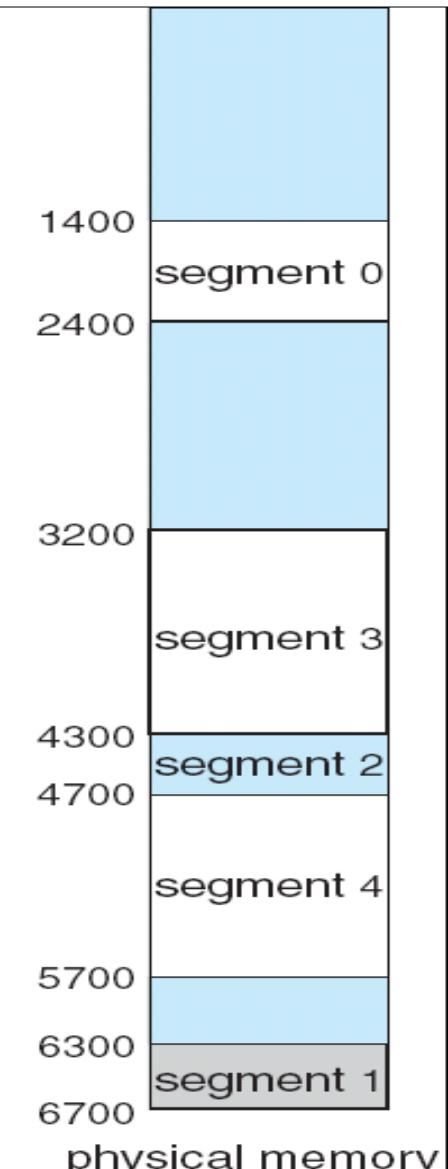
- E.g., 16-bit Intel 8086 & 8088
- Logical address consists of a two tuple:  
    <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
  - **base** – contains the starting physical address where the segments reside in memory
  - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program; segment number **s** is legal if **s < STLR**

# Example of Segmentation

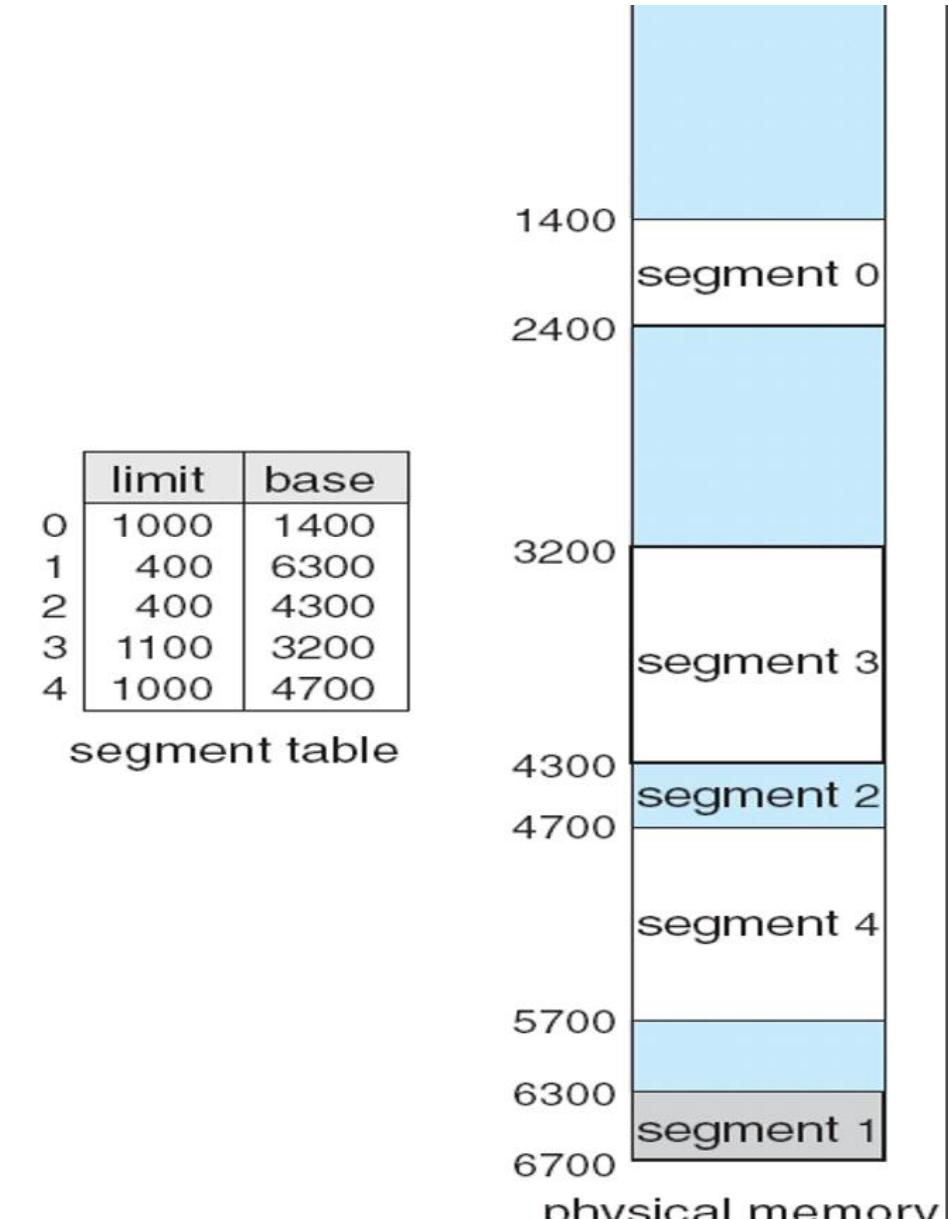
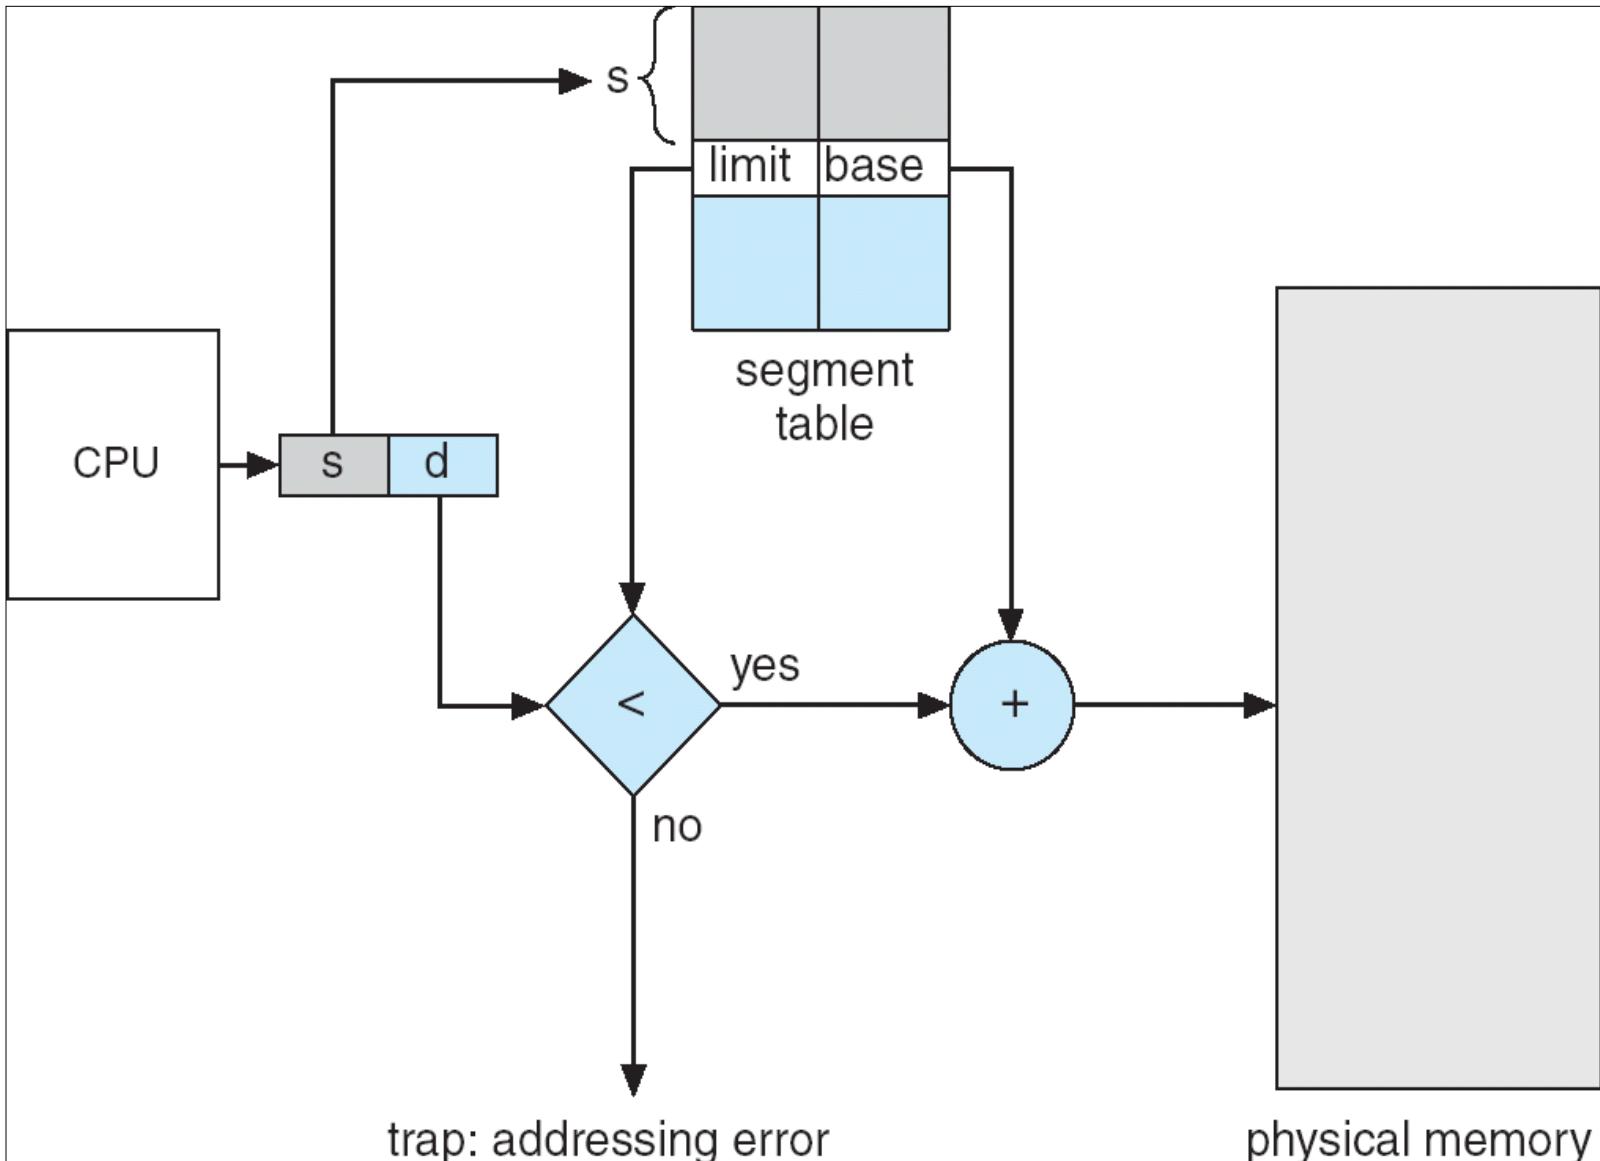


	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table



# Segmentation Hardware

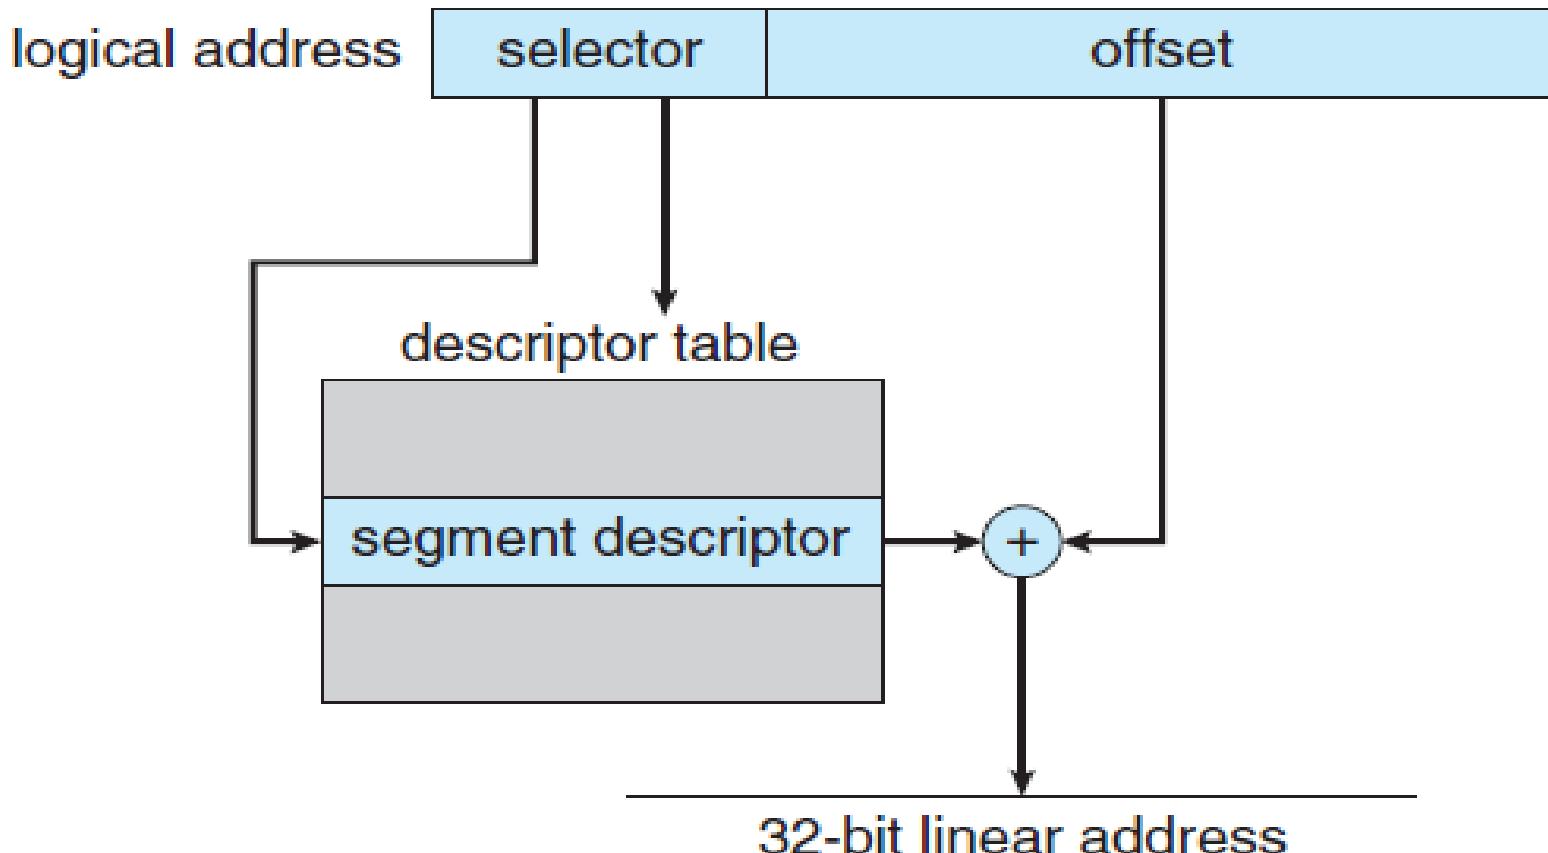
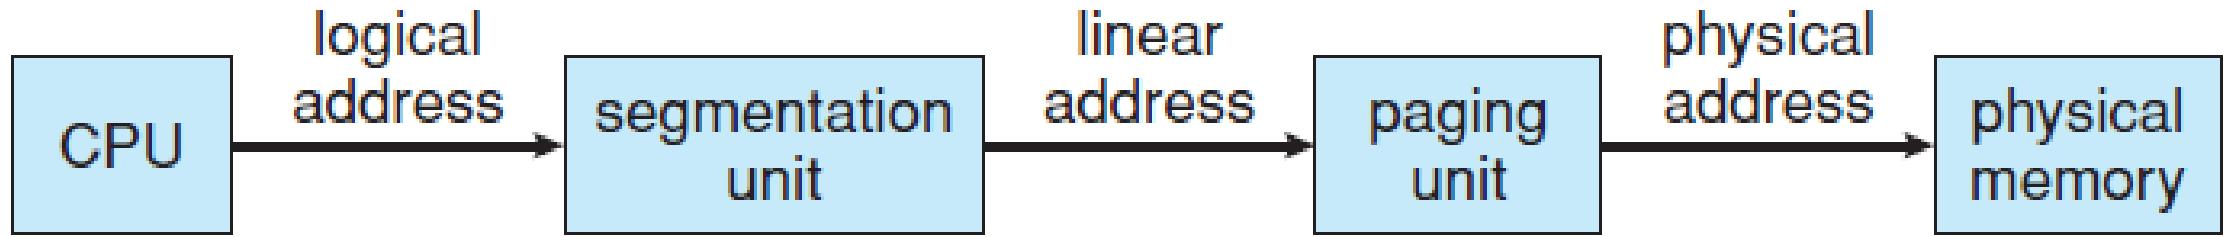


# Segmentation with paging

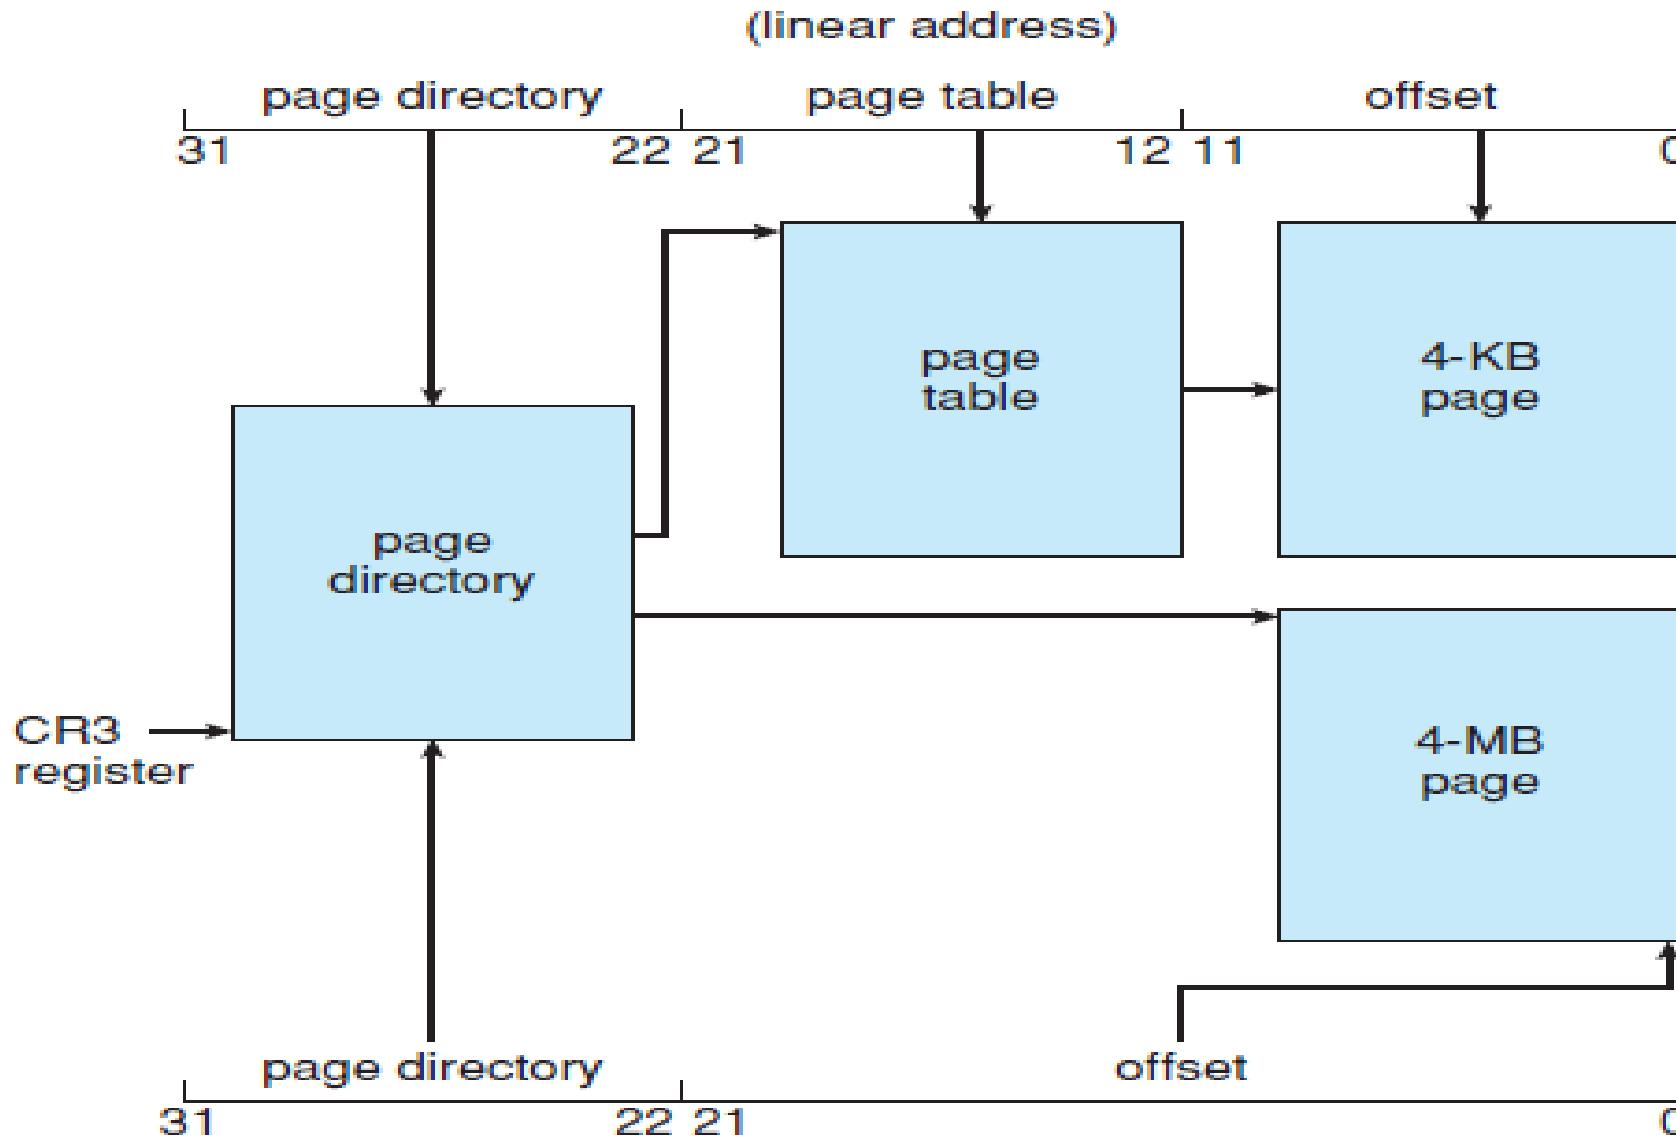
---

- Paging is transparent to the running process
- Segmentation requires that programs somehow be structured so that they are divided into logical parts with different address space
- With the proper hardware design a segmented program architecture can be run in combination with a paged memory architecture
  - Each segment is paged
- E.g., Intel 32-bit architecture (IA-32)

# Intel 32-bit architecture (IA-32)



# Intel 32-bit architecture (IA-32)

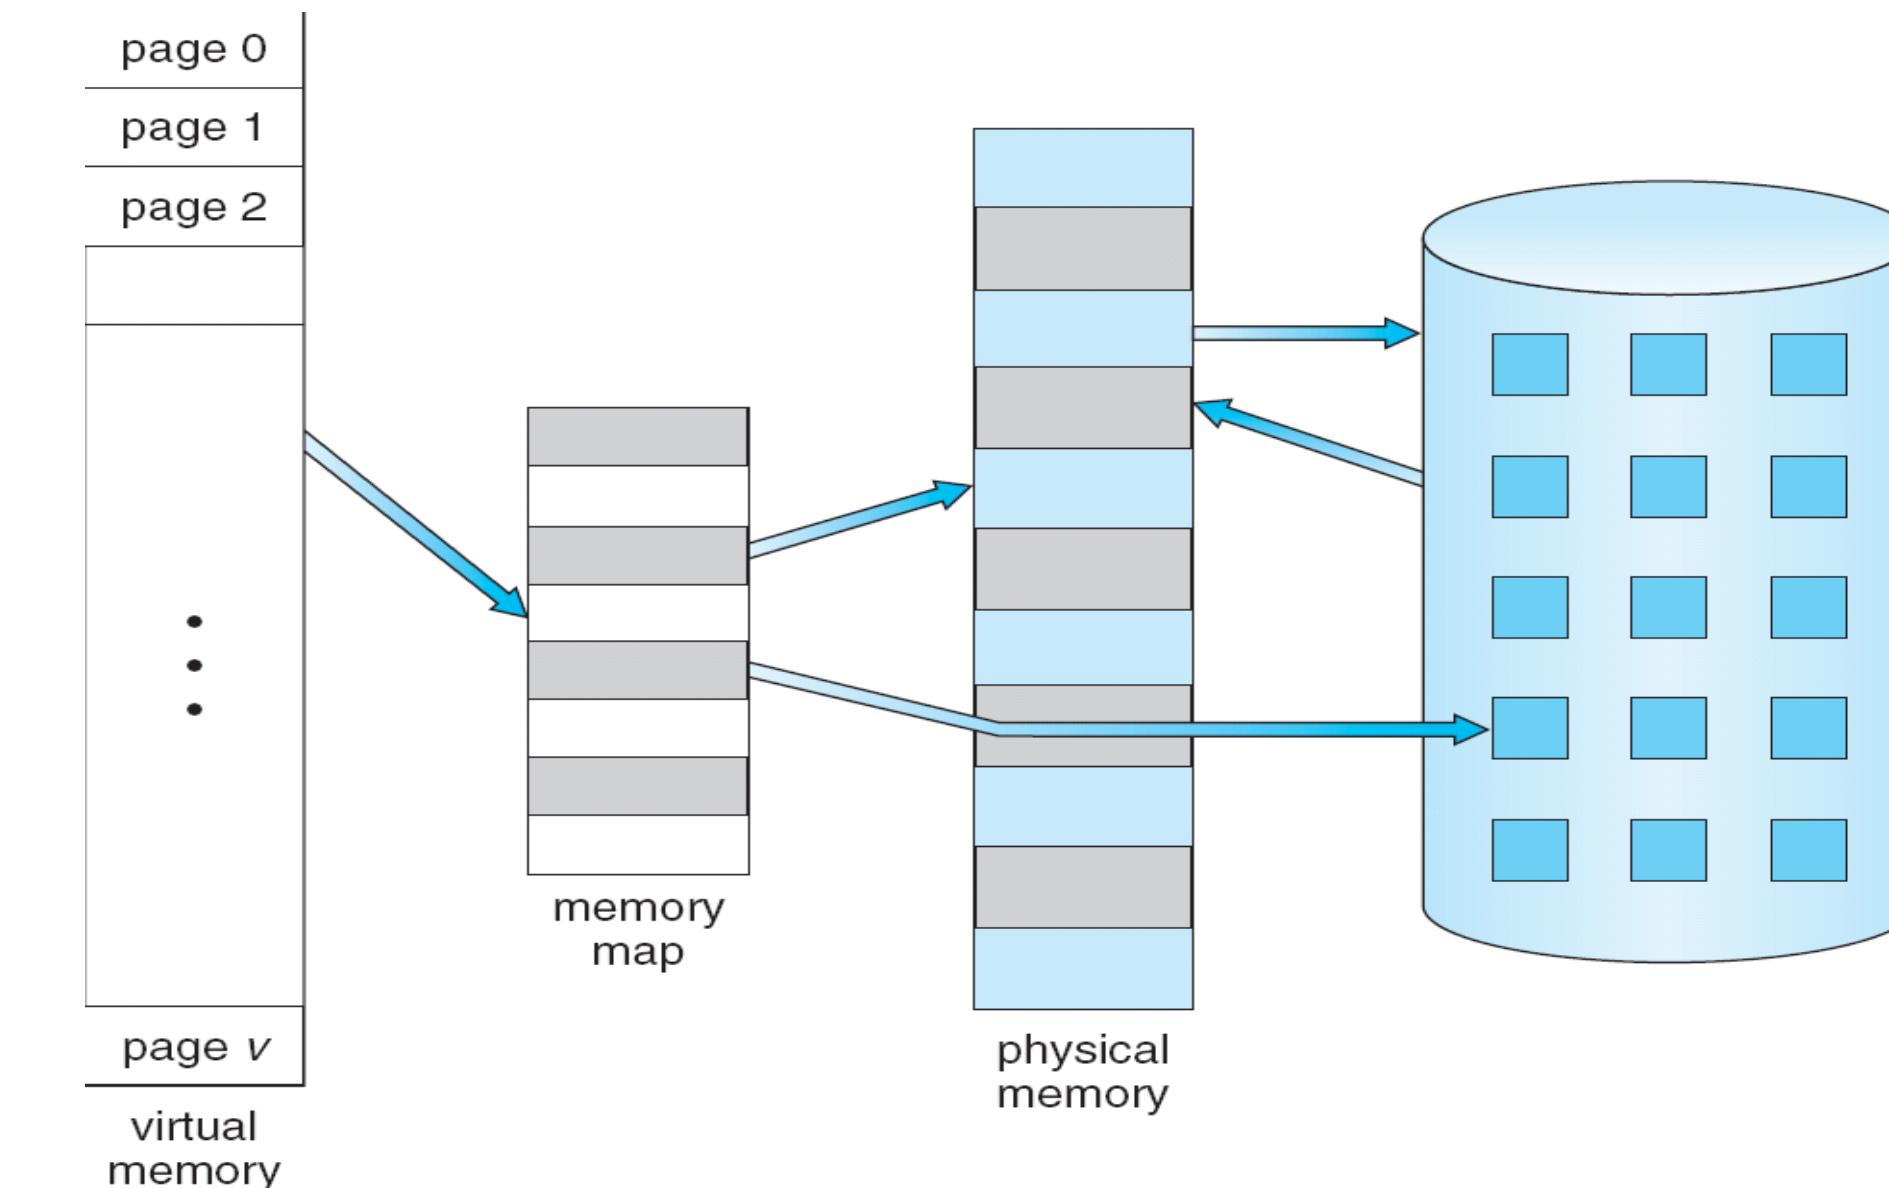


# Virtual Memory

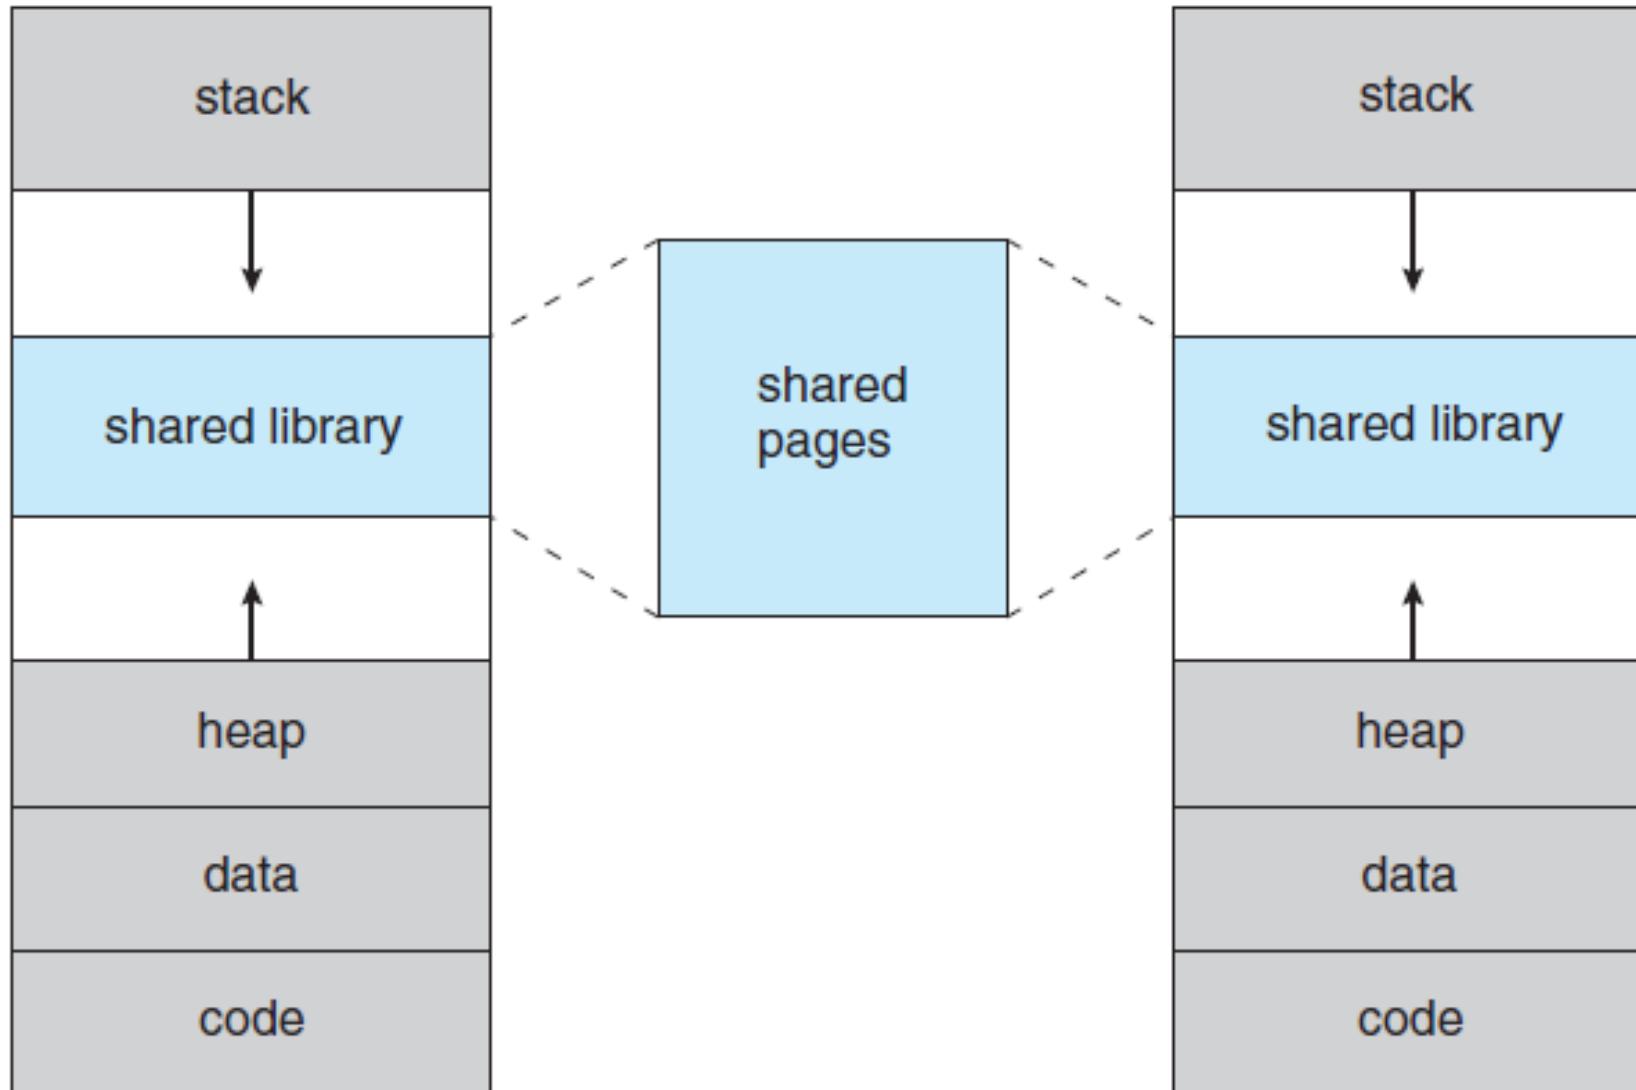
---

- Separation of user logical memory from physical memory
  - Only part of the program needs to be in memory for execution
  - Logical address space can therefore be much larger than physical address space
  - Allows address spaces to be shared by several processes
  - Allows for more efficient process creation
- Virtual memory can be implemented via:
  - Demand paging
  - Demand segmentation

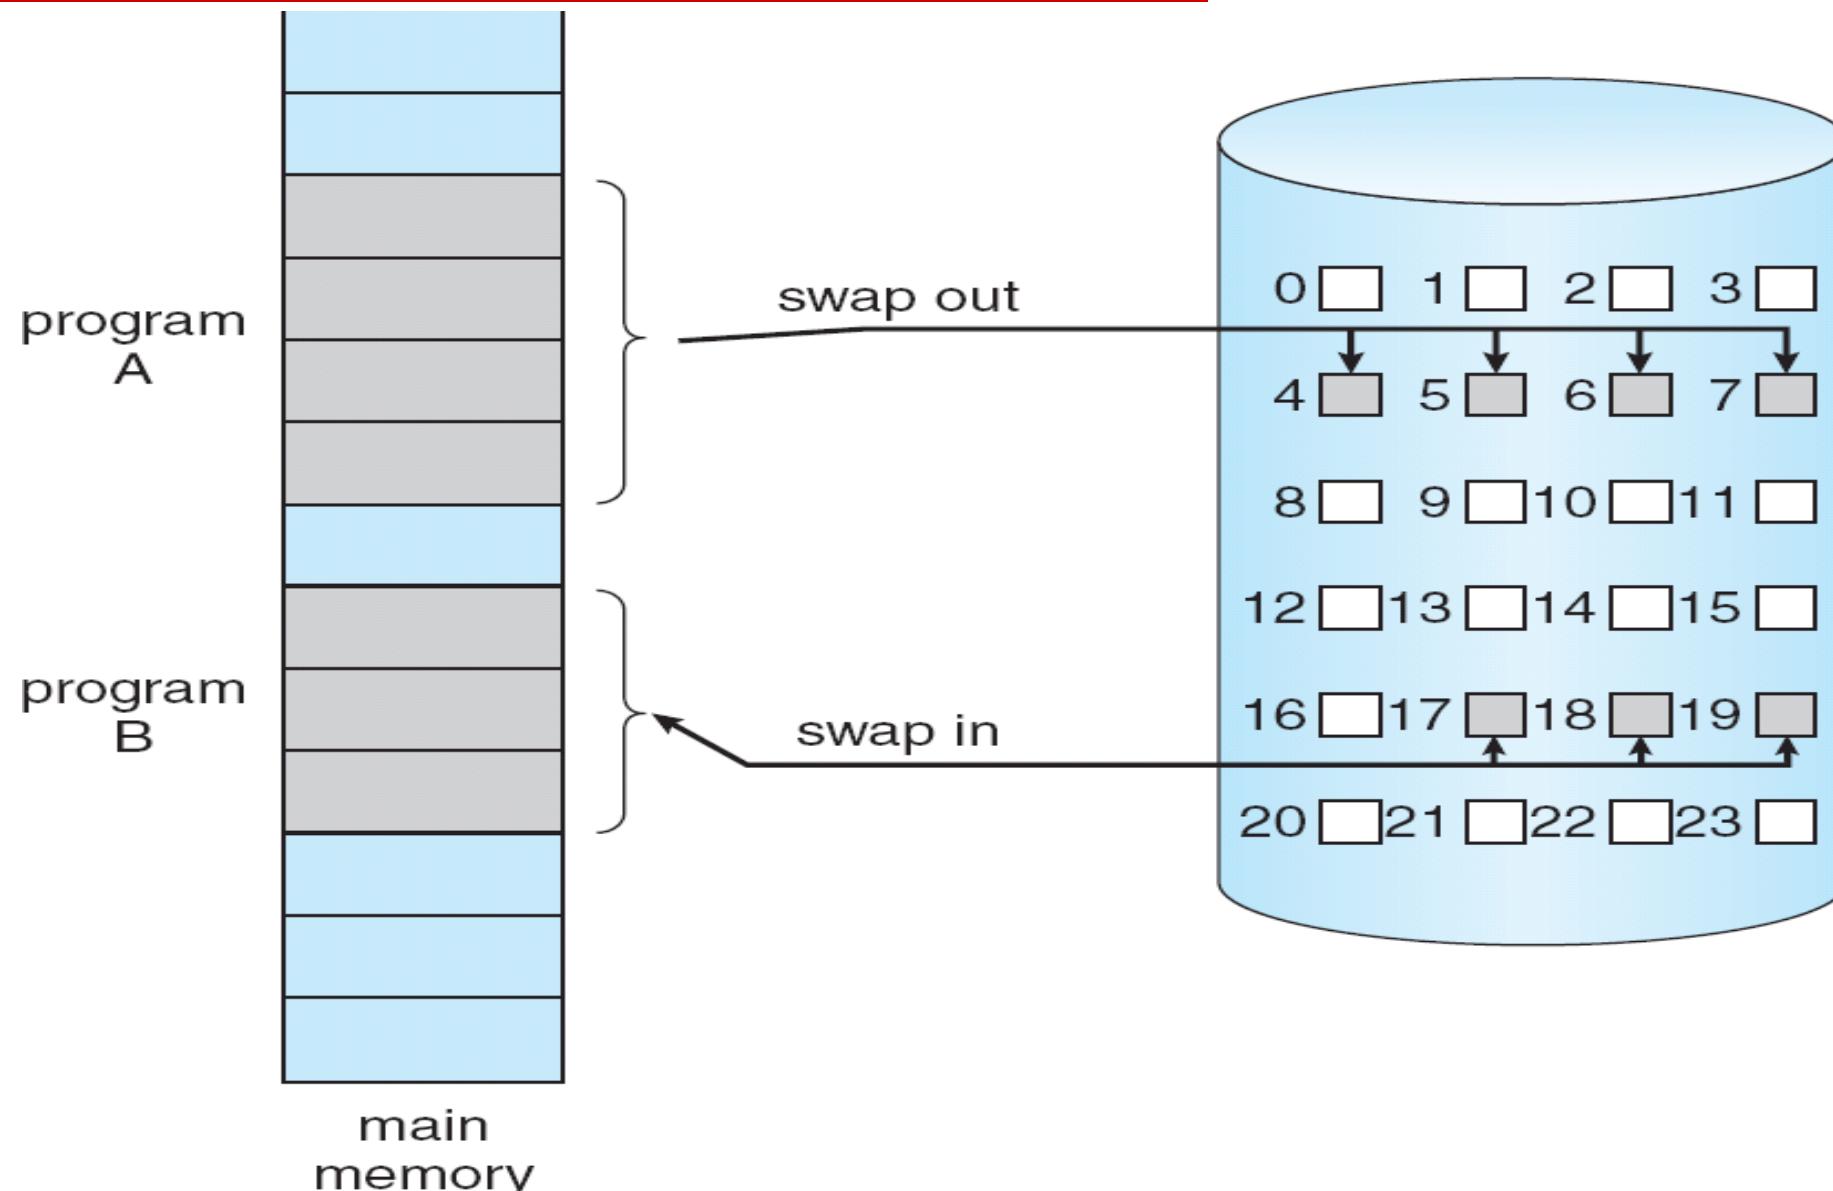
# Virtual Memory That is Larger Than Physical Memory



# Shared pages with virtual memory



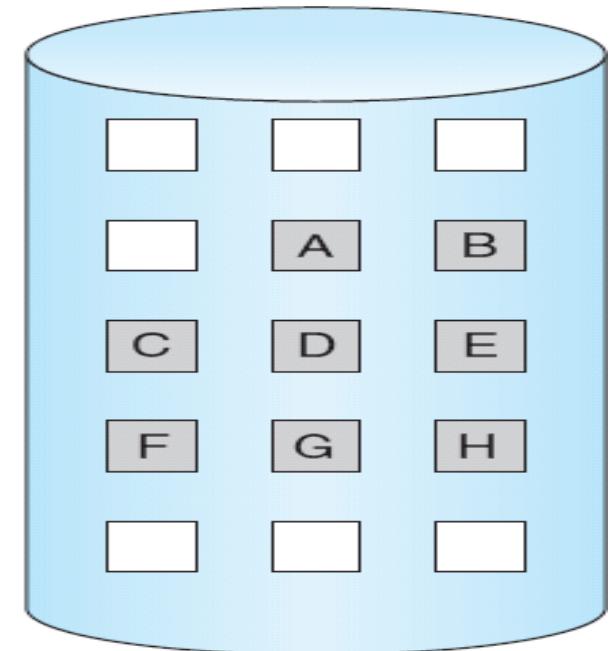
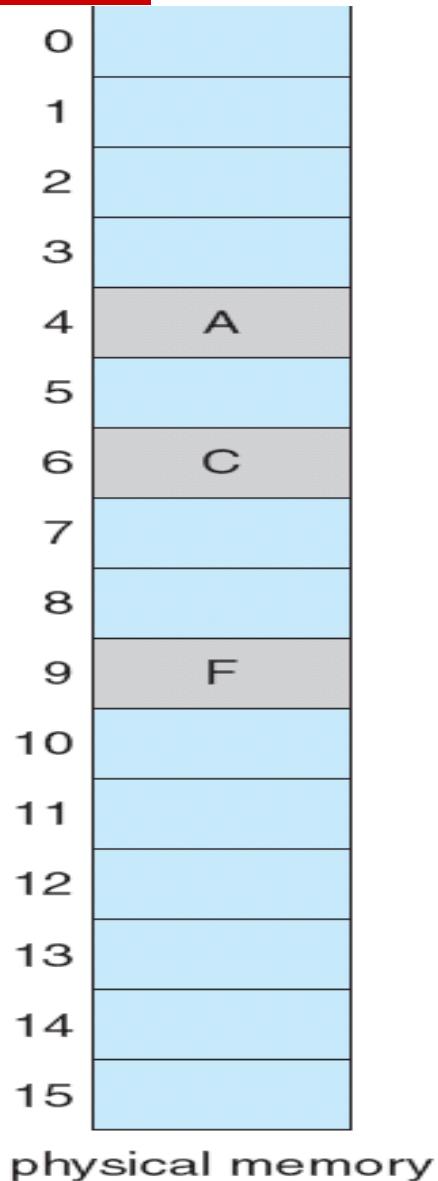
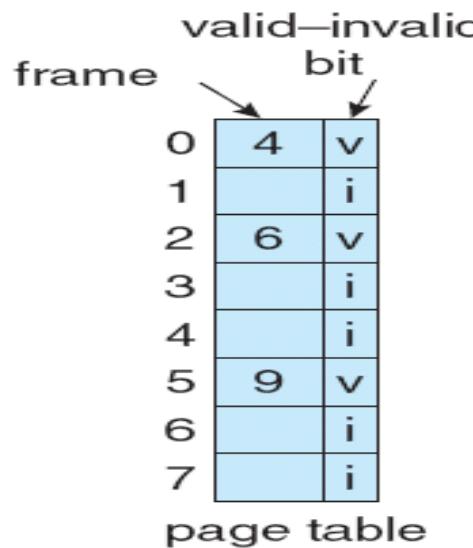
# Transfer of a Paged Memory to Contiguous Disk Space



# Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H

logical memory



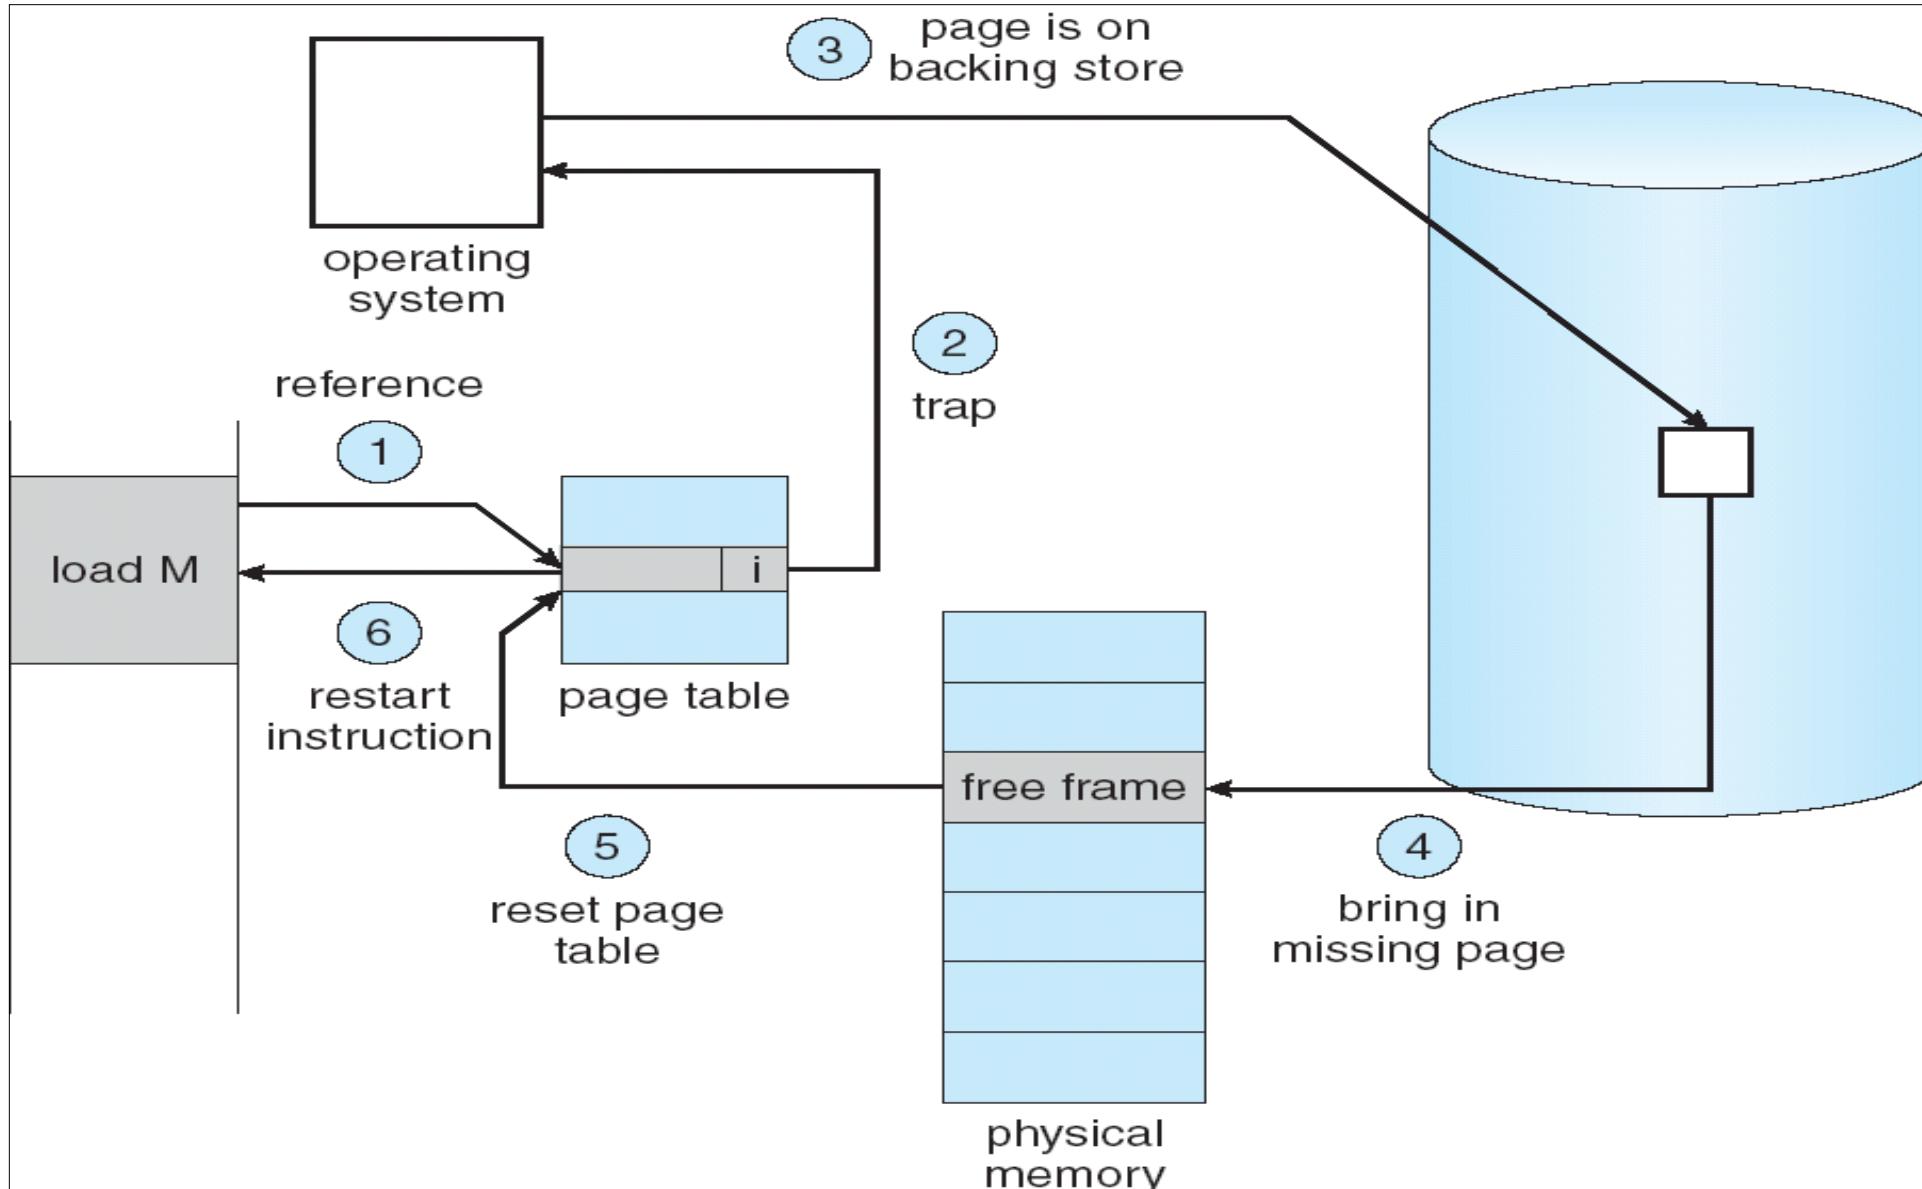
With each page table entry a valid-invalid bit is associated

**v** ⇒ in-memory

**i** ⇒ not-in-memory

Initially the bit is set to **i** on all entries

# Steps in Handling a Page Fault



# Performance of Demand Paging

---

- Page Fault Rate  $0 \leq p \leq 1.0$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault

- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead } ) \end{aligned}$$

# Demand Paging Example

---

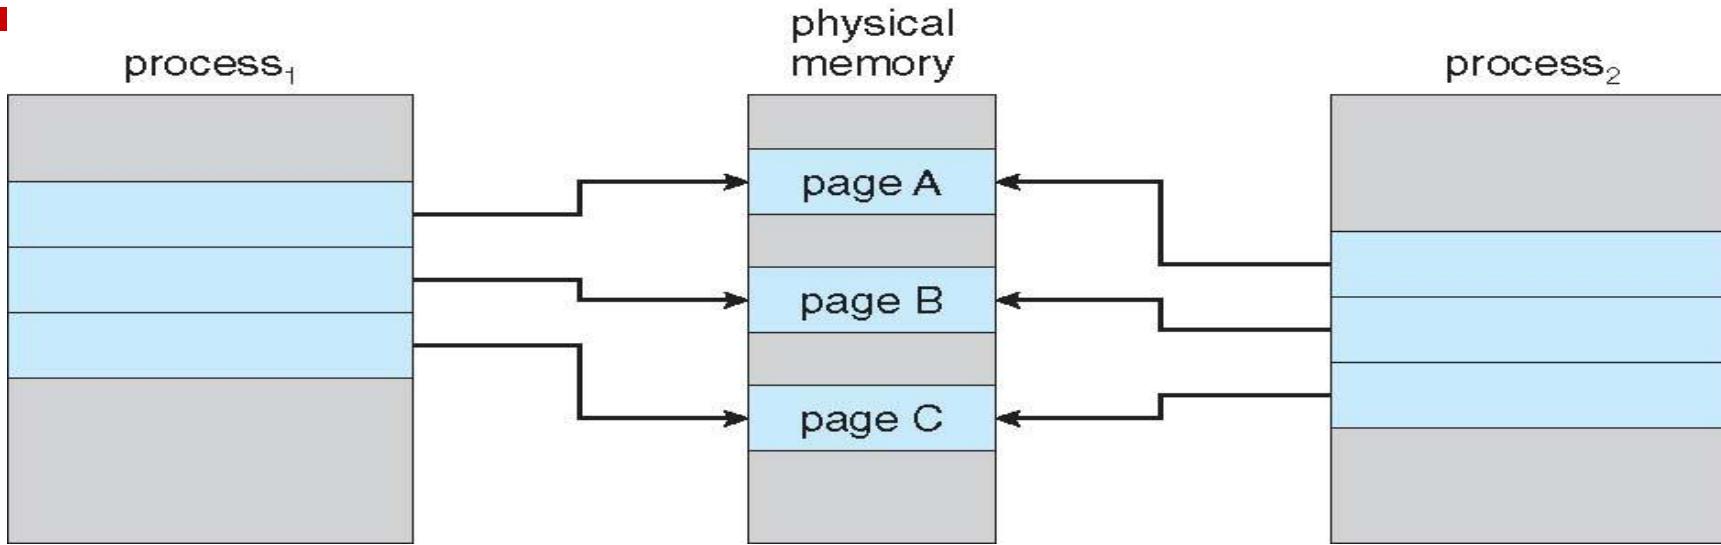
- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $$\begin{aligned} \text{EAT} &= (1 - p) \times 200 + p (8 \text{ milliseconds}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800 \end{aligned}$$
- If one access out of 1,000 causes a page fault, then  
 $\text{EAT} = 8.2 \text{ microseconds}$   
This is a slowdown by a factor of 40!!

# Copy-on-Write

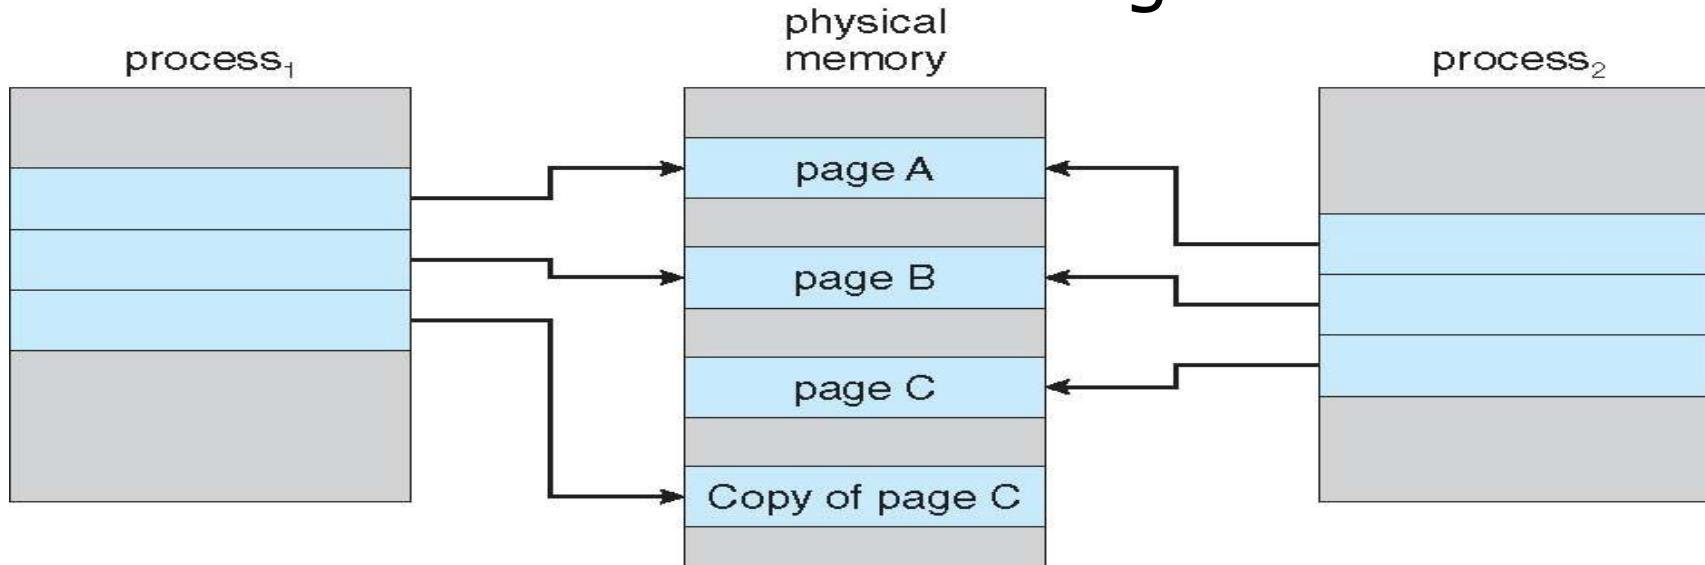
---

- **Copy-on-Write** (COW) allows both parent and child processes to initially **share** the same pages in memory
  - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied
- vfork() variation on fork() system call has parent suspend and child using copy-on-write address space of parent
  - Designed to have child call exec()
  - Very efficient

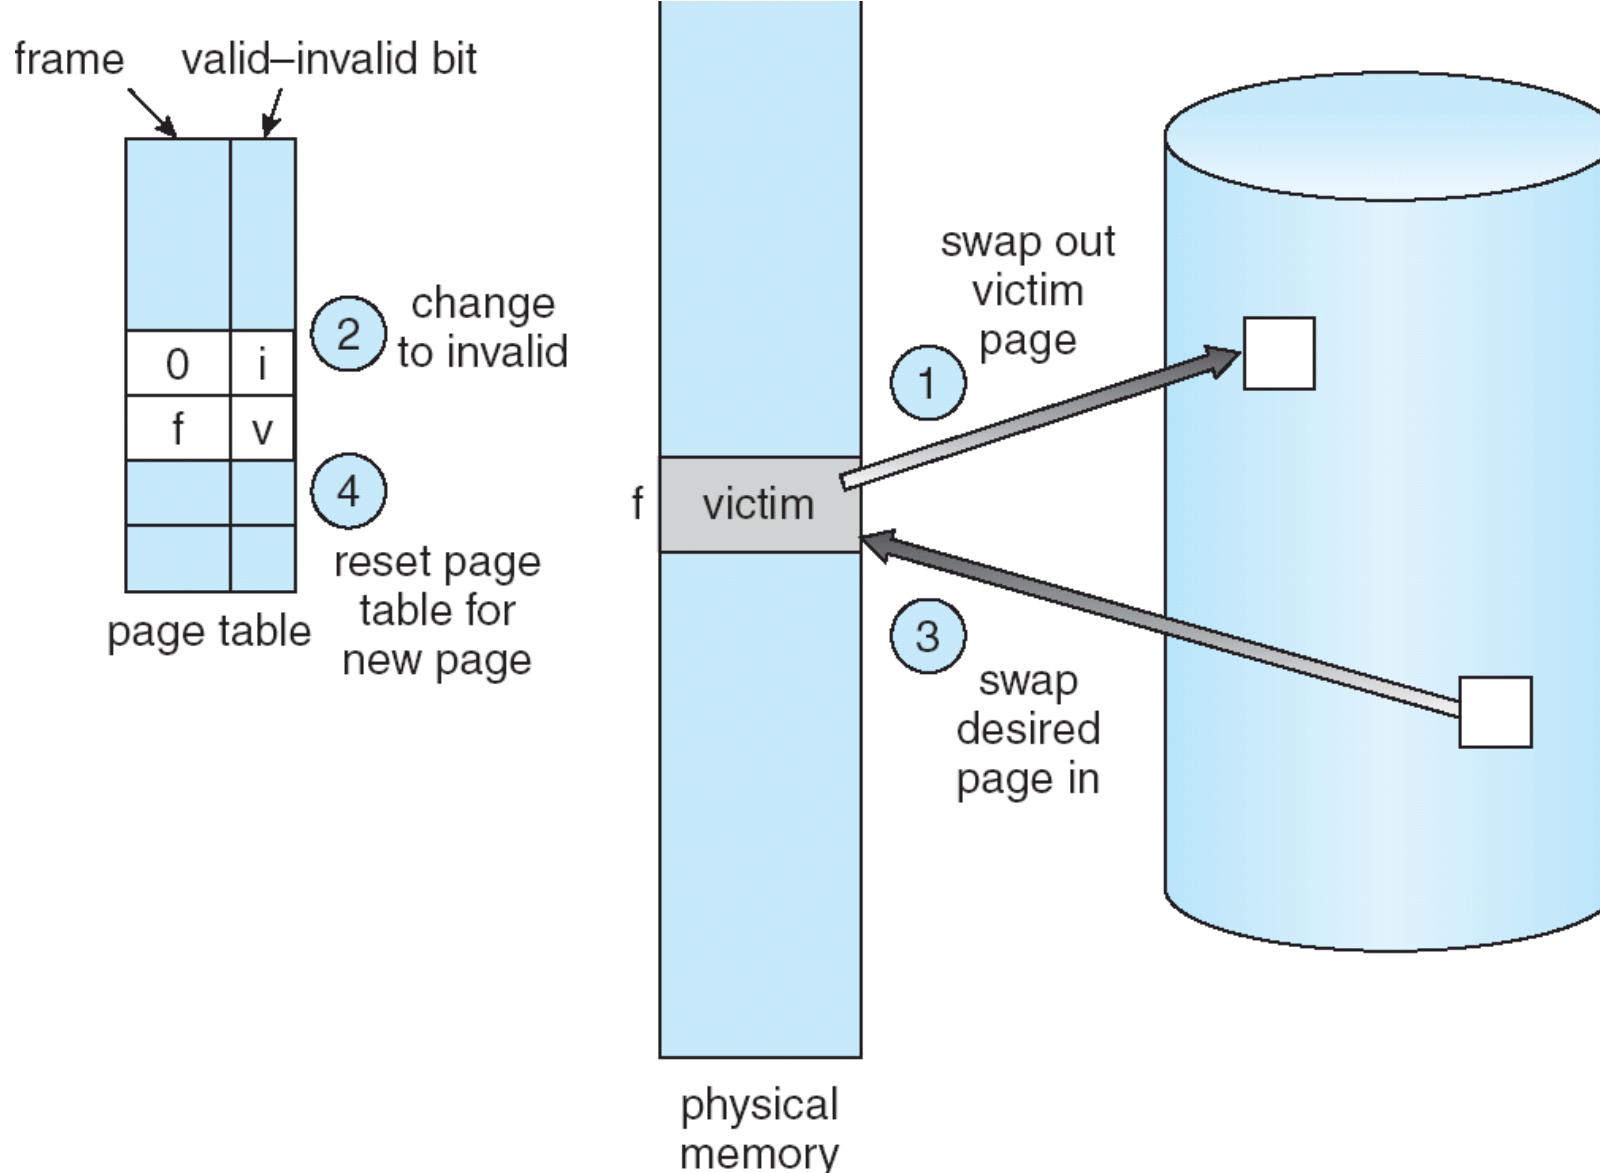
# Before Process 1 Modifies Page C



# After Process 1 Modifies Page C

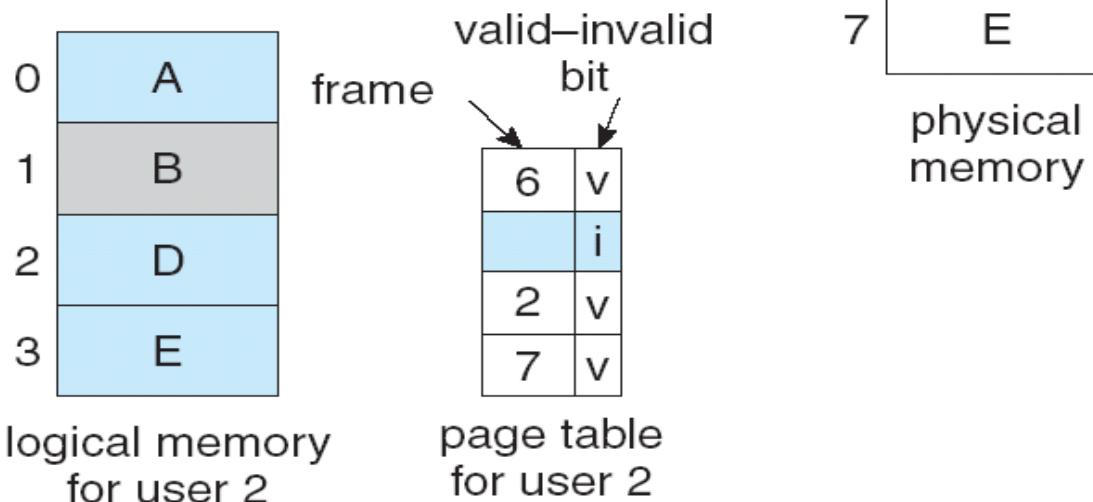
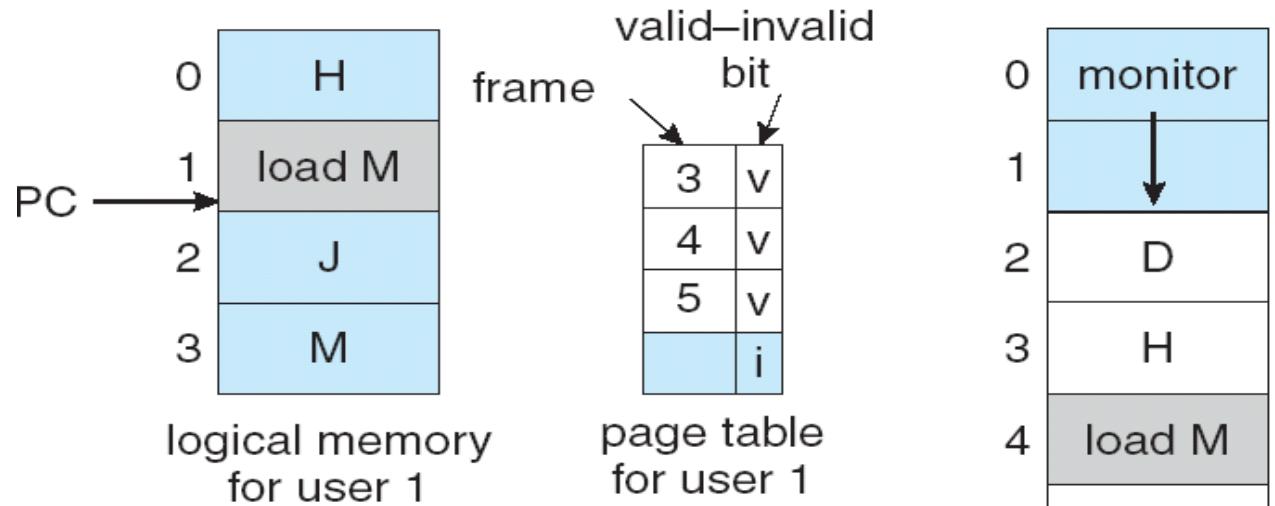


# What happens if there is no free frame?



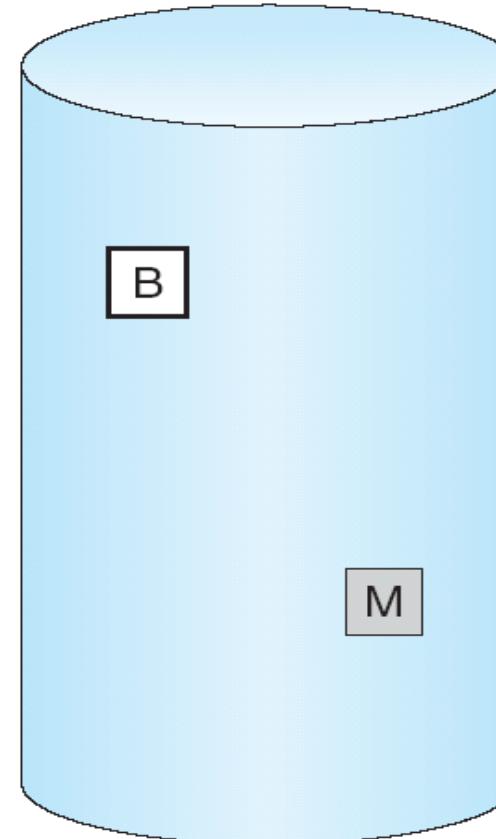
- Page replacement – find some page in memory, but not really in use, swap it out
  - algorithm
  - performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times

# Need For Page Replacement



physical memory

0	monitor
1	
2	D
3	H
4	load M
5	J
6	A
7	E



# Page Replacement Algorithms

---

- Want lowest page-fault rate
- Algorithms
  - First-in First-out (FIFO)
  - Optimal Page Replacement
  - Least Recently Used (LRU)
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- In all our examples, the reference string is

**1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5**

# First-In-First-Out (FIFO) Algorithm

---

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)

1	1	4	5
2	2	1	3
3	3	2	4

**9 page faults**

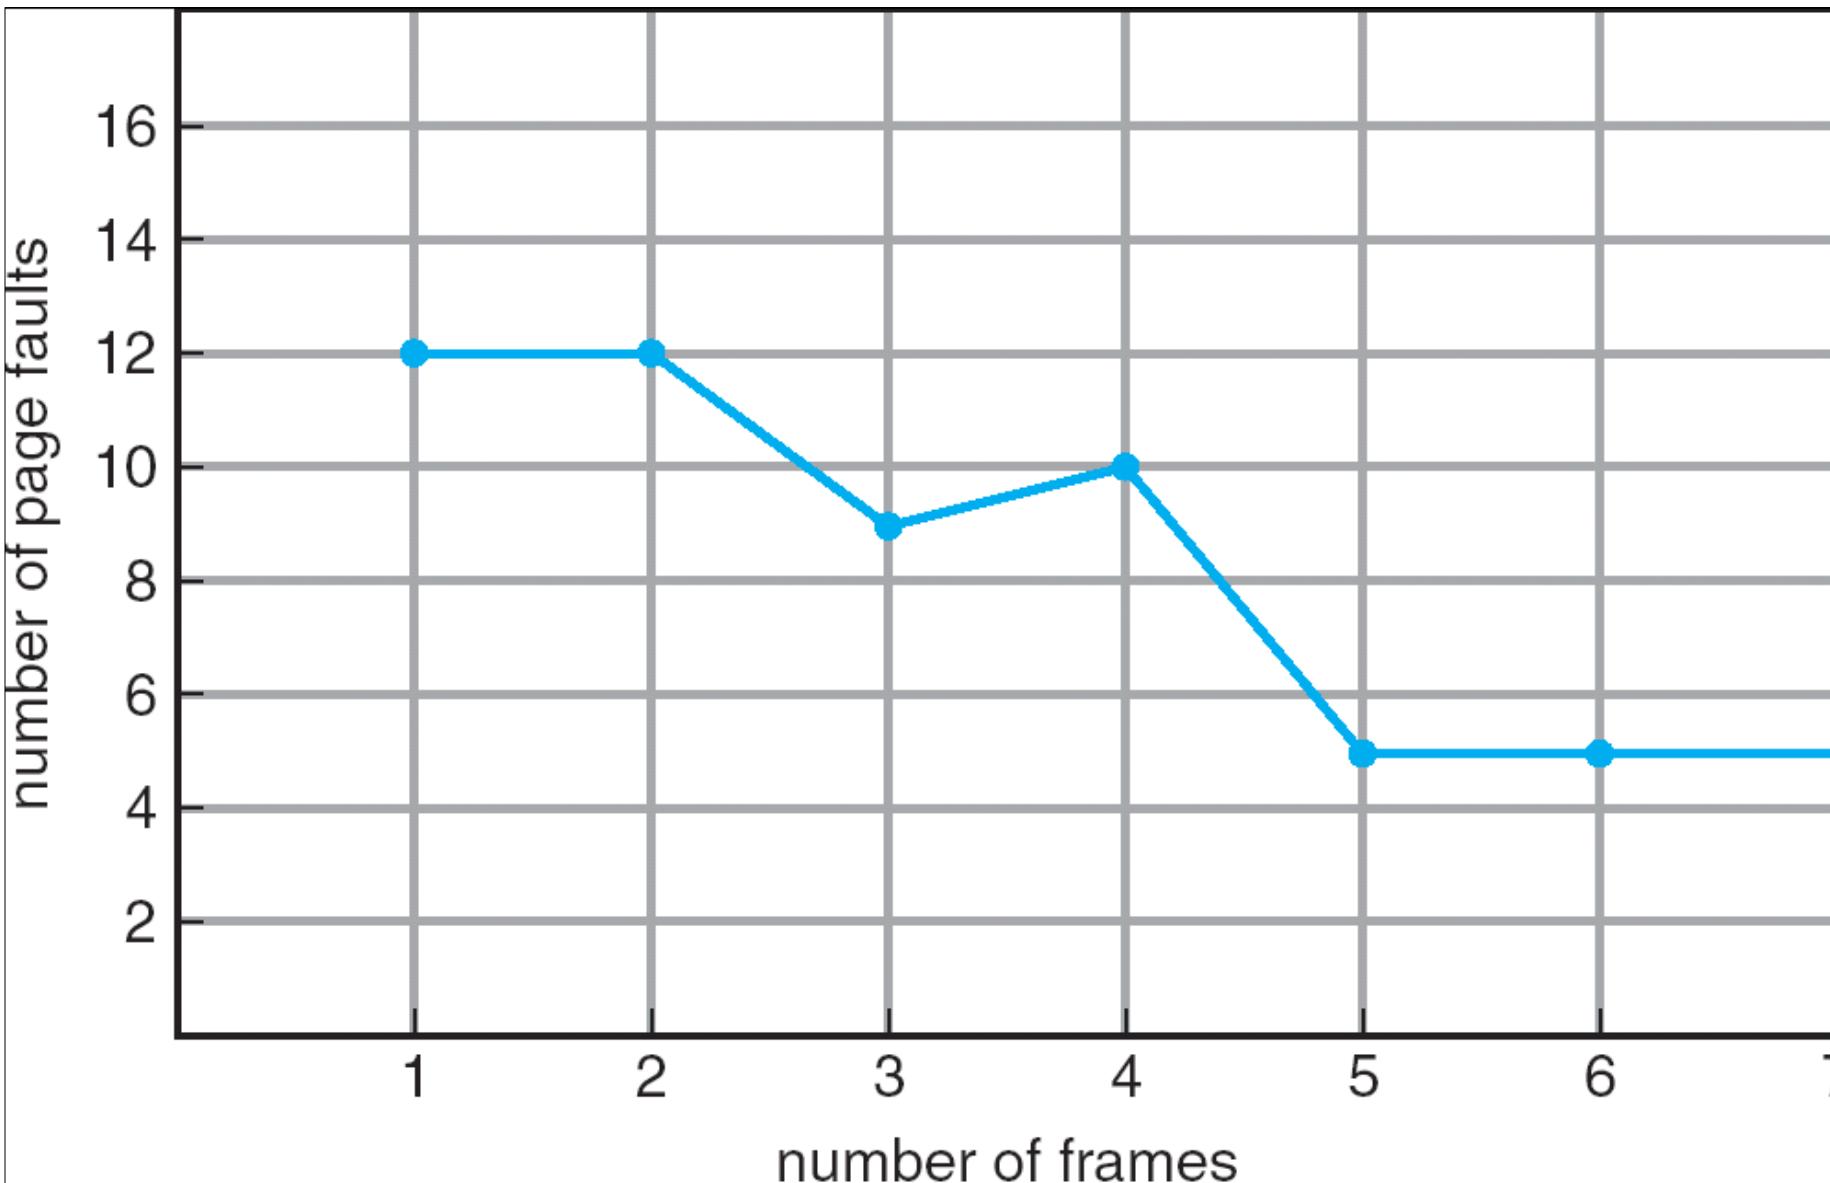
- 4 frames

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

**10 page faults**

Belady's Anomaly: more frames  $\Rightarrow$  more page faults

# FIFO Illustrating Belady's Anomaly

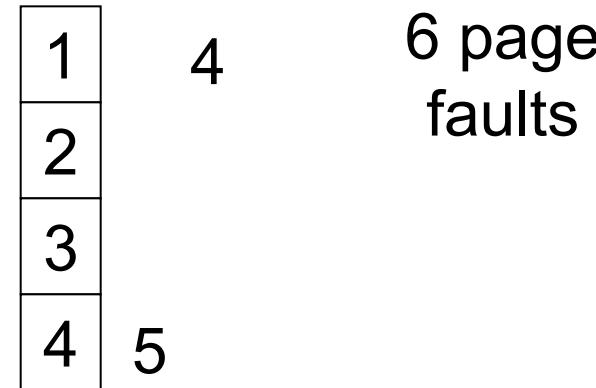


# Optimal Algorithm

---

- Replace page that will not be used for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



- How do you know this?
- Used for measuring how well your algorithm performs

# Least Recently Used (LRU) Algorithm

---

- Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	<b>5</b>	5	4	4
4	4	<b>3</b>	3	3

- Counter implementation
  - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
  - When a page needs to be changed, look at the counters to determine which are to change

# Global vs. Local Allocation

---

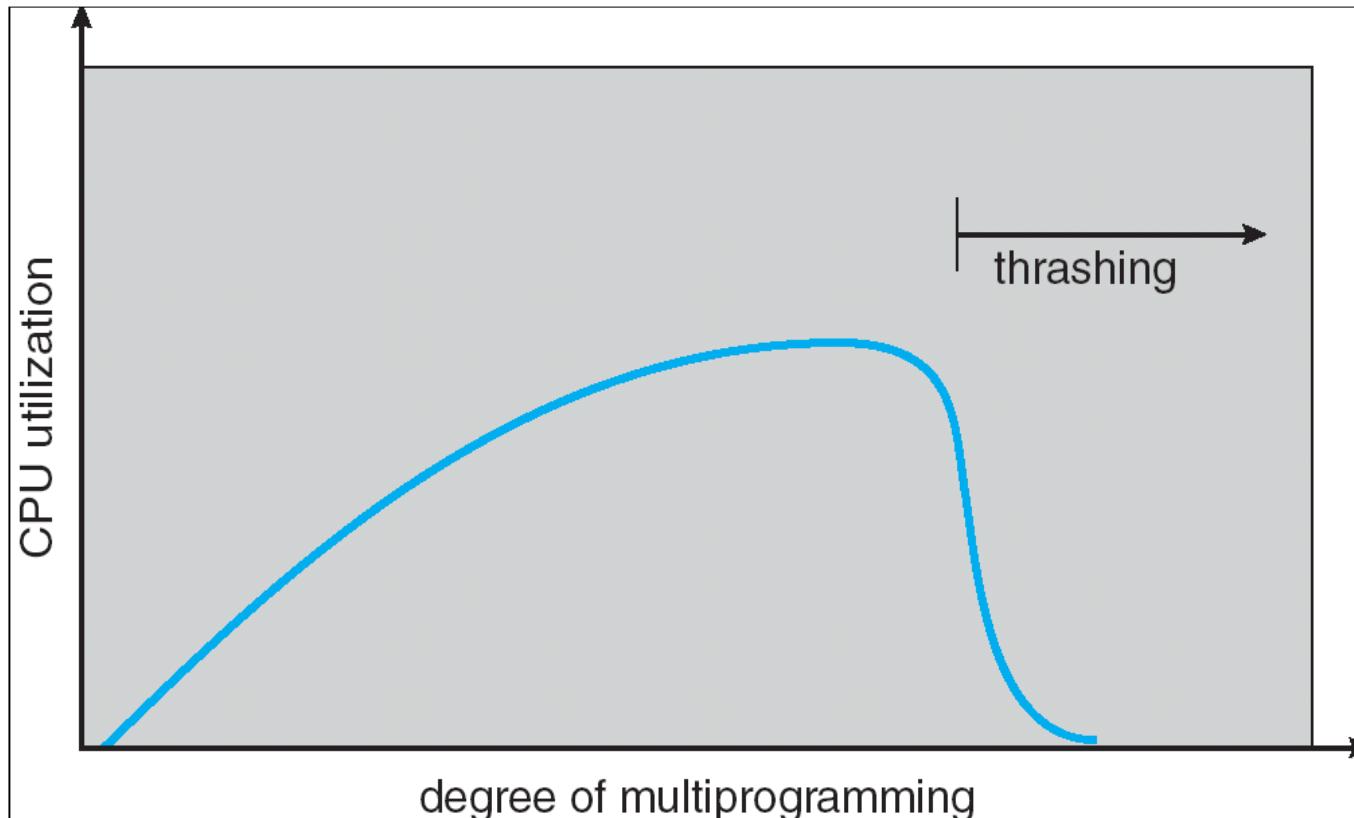
## Global replacement

- process selects a replacement frame from the set of all frames
- process can take a frame from another

## Local replacement

- each process selects from only its own set of allocated frames

# Thrashing



**Degree of multiprogramming = #processes**

- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
  - low CPU utilization
  - operating system thinks that it needs to increase the degree of multiprogramming
  - another process added to the system
- **Thrashing** ≡ a process is busy swapping pages in and out

# Kernel Memory Allocation

Treated differently from user memory

Often allocated from a free-memory pool

- requests memory for structures of varying sizes
- Some kernel memory needs to be contiguous
  - e.g., for device I/O

Two techniques

- Buddy
- Slab

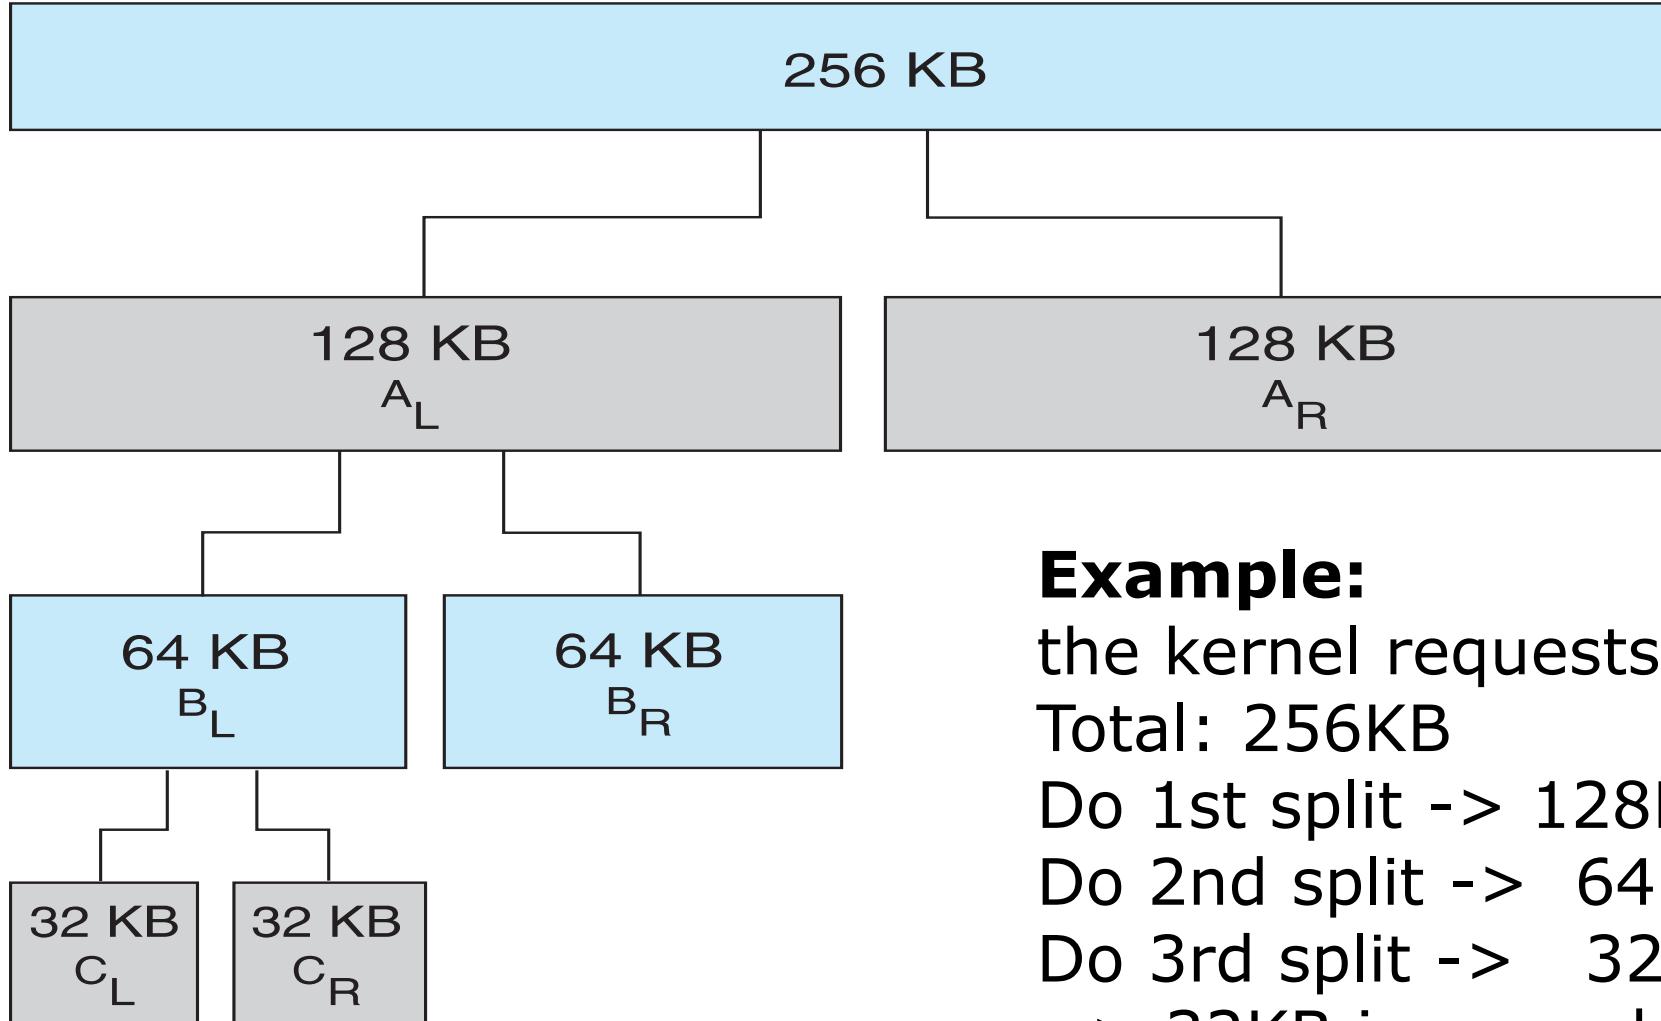
# Buddy System

---

- Allocates memory from fixed-size segment consisting of physically-contiguous pages
- Memory allocated using **power-of-2 allocator**
  - Satisfies requests in units sized as power of 2
  - Request rounded up to next highest power of 2
  - When smaller allocation than the one that is available is needed, the current chunk is split into two buddies of next-lower power of 2
  - Continue until appropriate sized chunk available
- Advantage – quickly **coalesce** unused chunks into larger chunk
- Disadvantage - fragmentation

# Buddy System Allocator

physically contiguous pages



## Example:

the kernel requests 21KB  
Total: 256KB  
Do 1st split -> 128KB  
Do 2nd split -> 64KB  
Do 3rd split -> 32KB  
=> 32KB is a good size

# Slab Allocator

---

- **Slab** is one or more physically contiguous pages
- **Cache** consists of one or more slabs
- Single cache for each unique kernel data structure
  - Each cache filled with **objects** – instantiations of the data structure
- If slab is full of used objects, next object allocated from empty slab
  - If no empty slab, a new slab is allocated
- Benefits include
  - no fragmentation
  - fast memory request satisfaction

# Slab Allocation

